Two inferential targets, different returns to case volume: theory and evidence for epidemic reproduction numbers

Abstract

This study asks why additional reported cases sharpen an estimate of the current reproduction number but do much less for transmission across periods. In a branching process with environmental variation, overdispersed offspring and binomial reporting, a conditionally unbiased estimate from m cases has mean squared error A/m about the current value R_t and $v_R + A/m$ about the mean across periods, because between-period variation of R_t is irreducible sampling variation for the second target, so precision levels off beyond the balance scale $m_\times = A/v_R$. Volatile reporting adds a term independent of m. With independent periods and linear costs, a fixed budget buys more from further periods once m exceeds $\sqrt{(c_0/c_1)m_\times}$. In US influenza admissions during four early-growth seasons, an affine loss model gave an apparent balance ratio of 37.4 admissions per week (95% CI 19.1 to 71.0), though a power law fitted about as well. Binomial thinning steepened the slope at 0.41 (95% CI 0.31 to 0.53), near the Poisson value 0.45. State intercepts gave 17.3 (95% CI 3.3 to 51.4), not distinguishable from the cross-sectional fit. Simulated calibration gave an environmental variance of 0.040 (95% CI 0.026 to 0.056) at a 3.6-day mean generation interval and 0.013 to 0.026 at 2.2 to 2.9 days. Between 41% and 51% of the intercept was shared by all states in a week; 55% of state-weeks lay above the apparent ratio and 44% above the latent scale.

Keywords:
Reproduction number
Environmental stochasticity
Surveillance design
Reporting heterogeneity
Influenza hospitalisation
Diminishing returns

1. Introduction

Estimates of the time-varying reproduction number have become a routine output of infectious disease surveillance. They were used to track the transmission of Ebola virus disease in West Africa (Nouvellet et al., 2018), to assess non-pharmaceutical interventions against COVID-19 (Flaxman et al., 2020) and to describe the early spread of SARS-CoV-2 (Kucharski et al., 2020). They are now produced routinely for many countries (Huisman et al., 2022), with practical guidance on their estimation (Gostic et al., 2020; Steyn et al., 2025). Their value for situational awareness depends on how precisely they can be estimated from the data that surveillance systems actually collect.

The instantaneous reproduction number, R_t, is the expected number of secondary infections generated by an individual infected at time t if conditions remained as they were at that time (Fraser, 2007). It is usually estimated from incidence through the renewal equation, which weights past incidence by the generation interval (Wallinga and Teunis, 2004; Cori et al., 2013; Thompson et al., 2019). The generation interval is the time between infection of a primary case and infection of its secondary cases. For a given generation interval, the posterior uncertainty of such estimates shrinks as the number of cases in the estimation window increases (Cori et al., 2013). Parag and colleagues improved estimation at low incidence (Parag, 2021). They also quantified in closed form the information that incidence carries about R_t, including the loss caused by time-varying reporting fractions (Parag et al., 2022), and derived the limits this information places on timely interventions (Parag et al., 2025). Extensions handle temporally aggregated counts (Nash et al., 2023; Ogi-Gittins et al., 2024) and the difference between the reproduction number of infections and the apparent reproduction number of a delayed observation series (Eales and Riley, 2024).

These precision results measure error against the realised value of R_t in the period being estimated, and that error vanishes as case counts grow. Many uses of R_t estimates refer instead to a quantity that persists across periods, such as the typical level of transmission in a season or the value expected next week. Policy evaluations likewise average over weeks with different climatic conditions, school calendars and contact patterns, all of which modulate influenza transmission (Cauchemez et al., 2008; Shaman and Kohn, 2009; Shaman et al., 2010). Between-period variation in R_t is shared by all chains of transmission in a population at a given time. Population ecology distinguishes such environmental stochasticity from demographic stochasticity, which arises from independent individual outcomes (Engen et al., 1998; Rohani et al., 2002; Lande et al., 2003), and mechanistic epidemic models have included both forms (Bretó et al., 2009). Demographic variance per individual falls with population size, whereas environmental variance does not, and the population size at which the two contributions are equal is a standard summary of extinction risk (Lande, 1993; Lande et al., 2003).

Overdispersed offspring distributions (Lloyd-Smith et al., 2005; Endo et al., 2020) and incomplete reporting add to the demographic side of this partition.

Surveillance data add a third source of variation. The proportion of infections that appear in a count changes over time with testing practice, care seeking and the completeness of facility reporting (Azmon et al., 2014; Stoner et al., 2019; Parag and Donnelly, 2022). In the United States, hospitals were required to report influenza admissions to the National Healthcare Safety Network (NHSN) from February 2022 (Centers for Disease Control and Prevention, 2022; Frutos et al., 2026). Reporting was voluntary from May to October 2024 and required again from November 2024 (Centers for Medicare and Medicaid Services, 2024; Centers for Disease Control and Prevention, 2026). Hospital admission series of this kind are increasingly used for forecasting and real-time assessment (Mathis et al., 2024). Variation in reporting completeness that persists within a week acts on every case in that week and therefore behaves like environmental variation, not like sampling noise.

Sample size calculations for surveillance usually assume that more cases give proportionally more precision (Cochran, 1977), and analyses of reproduction number estimation have concentrated on the informativeness of incidence for period-specific values (Parag et al., 2022; Bhatt et al., 2023). A variance that is the sum of a term shared within a unit and a term that falls with the number of observations in the unit is, however, familiar from several fields. It is the variance of a cluster mean in two-stage sampling (Cochran, 1977), and it appears in sample size calculations for cluster-randomised trials. There the between-cluster coefficient of variation sets the number of clusters needed and the cost ratio sets the optimal cluster size (Hayes and Bennett, 1999; Raudenbush, 1997). Designs for detecting trends in repeated ecological surveys face the same structure, because variation shared by all sites in a year limits what additional sites can achieve (Urquhart and Kincaid, 1999; Kincaid et al., 2004). State-space models of population dynamics show that process and observation variance are often difficult to separate from a single series (Knape, 2008; Auger-Méthé et al., 2016). Stochastic epidemic models fitted to incidence include environmental noise in transmission alongside demographic noise (Bretó et al., 2009; He et al., 2010; Dureau et al., 2013; King et al., 2015). The form $v_R + A/m$ is therefore classical. What surveillance for the reproduction number adds is the content of each term. The term that falls with case volume depends on the mean reproduction number through overdispersed offspring and binomial reporting, volatile reporting completeness adds a term that does not fall with case volume, and a budget is spread over several periods.

An empirical test also faces a measurement problem. Weekly growth statistics computed from admissions are apparent quantities whose variance depends on the delay from infection to admission and on how each statistic aggregates growth, so their intercepts cannot be read directly as transmission variance. Observations from the same calendar week share national influences, and interval estimates that treat them as independent are too narrow (Owen, 2007; Cameron et al., 2011).

The central point is a distinction between inferential targets. For the current value R_t, the error of a case-count estimate is counting error, and it vanishes as cases per period accumulate. For the mean reproduction number across periods, the true variation of R_t from one period to the next is part of the error. It is not counting error that more cases per period would remove, nor is it uninformative noise. It is irreducible sampling variation of the cross-period target, and only more periods reduce it. Our contribution is not the variance partition itself. It is a mechanistic version for case-count surveillance under negative binomial offspring and binomial reporting, the balance scale $m_\times$ at which its two parts are equal, its consequences for volatile reporting and for dividing a fixed budget between cases per period and number of periods, and an empirical calibration. For the calibration we fit the relative loss of two apparent weekly statistics in NHSN influenza admissions against case volume and map their intercepts to the latent environmental variance through a data-matched simulation, keeping the apparent scale of each statistic distinct from the latent scale of the model (Table 1). Two further analyses address what a cross-sectional fit cannot: binomial thinning of the observed counts tests whether the term that falls with case volume scales as the model predicts, and state intercepts separate the cross-sectional slope from the slope within states. Finally, we decompose the intercept into national-week and state-specific components and compare the mandatory and voluntary reporting regimes.

2. Model and theory

2.1. Transmission, reporting and two estimation targets

Consider transmission over one generation, indexed by t, in which n infected individuals act as primary cases. Conditional on the instantaneous reproduction number R_t, the number of secondary infections caused by primary case i, $X_{i,t}$, is assumed to follow a negative binomial distribution, which captures individual heterogeneity in transmission (Lloyd-Smith et al., 2005). The offspring distribution is given by:

$$X_{i,t} \mid R_t \sim \mathrm{NB}\left(k,\ \frac{k}{k+R_t}\right), \qquad \mathbb{E}(X_{i,t} \mid R_t) = R_t, \qquad \mathrm{Var}(X_{i,t} \mid R_t) = R_t + \frac{{R_t}^2}{k}, \tag{1}$$

where $k > 0$ is the dispersion parameter, smaller values of k correspond to greater heterogeneity, and $k \to \infty$ recovers Poisson offspring. Each infection is then reported independently with probability $\rho \in (0, 1]$. Because a sum of independent negative binomial variables with a common success probability is negative binomial, and because binomial thinning

Table 1. Notation and scale of the main quantities. Latent quantities belong to the transmission model and refer to one generation; apparent quantities are defined on the scale of a weekly statistic.

Symbol	Meaning	Scale	Values in this study
R_t; $\bar{R}$; v_R	Instantaneous reproduction number in period t; its mean and variance across periods	Latent, per generation	Baseline $\bar{R} = 1.15$, $v_R = 0.04$
m	Expected reported cases per period when $R_t = 1$ (case volume)	Latent, cases per generation	–
A; $m_\times = A/v_R$	Sampling coefficient and balance scale (Proposition 1)	Latent, cases per generation	Baseline 53.1; 30.3 per generation (95% CI 21.7 to 46.3) implied by the data, or 58.9 per week (95% CI 42.3 to 90.0)
τ; v_H; $m_{\times,H}$	Relative reporting volatility, variance it adds, generalised balance scale (Proposition 2)	Latent	$\tau = 0.15$ gives $v_H = 0.0307$ and $m_{\times,H} = 30.4$
m^*	Budget-optimal case volume per period with independent periods and linear costs (Proposition 3)	Same as the balance scale used	–
G_t; $R_{A,t}$	Weighted growth ratio and renewal-type ratio (Eqs. (13) and (14))	Apparent, weekly	–
D_t	Reporting scale, a weighted sum of recent weekly admissions	Admissions per week	–
a; b	Intercept and slope of the relative loss (Eq. (16))	Apparent, scale of each statistic	$a = 0.0150$ for R_A, 0.0840 for G
b/a	Apparent balance ratio	Admissions per week on the scale of each statistic	37.4 for R_A, 23.5 for G
$b_{\mathrm{w}}/\bar{a}_s$	Within-state apparent balance ratio, with state intercepts a_s	Admissions per week, scale of R_A	17.3 (95% CI 3.3 to 51.4)
c; b_P	Rate at which the slope rises under binomial thinning; its Poisson value	Scale of R_A	$c = 0.41$ (95% CI 0.31 to 0.53); $b_P = 0.45$
κ	Scale factor $\mathbb{E}(\hat{a})/(v_R/\bar{R}^2)$ from the calibration	Ratio	0.516 for R_A, 2.729 for G
v_R^{impl}	Latent environmental variance implied by the two-step inversion	Latent, per generation	0.040 for R_A, 0.043 for G at a mean generation interval of 3.6 days; 0.013 to 0.026 for R_A at 2.2 to 2.9 days

preserves this family, the reported count C_t satisfies:

$$C_t \mid n, R_t \sim \mathrm{NB}\left(nk,\ \frac{k}{k+\rho R_t}\right), \qquad \mathbb{E}(C_t \mid n, R_t) = \rho n R_t, \qquad \mathrm{Var}(C_t \mid n, R_t) = \rho n R_t + \frac{\rho^2 n {R_t}^2}{k}, \tag{2}$$

where $m = \rho n$ is the expected number of reported cases per period when $R_t = 1$. We refer to m as the case volume.

The model rests on four assumptions. Each infection is reported independently with a probability ρ that does not change over time (H1). Primary cases transmit independently and identically according to Eq. (1) (H2). The dispersion k does not depend on R_t or n (H3). The number of primary cases n is fixed by design and is independent of R_t (H4). Assumption H4 holds when n is set by contact tracing; in routine surveillance n is itself the outcome of earlier transmission, and the renewal formulation in Section 3.3 relaxes it.

Two scales recur in what follows (Table 1). The quantities of the model, namely R_t, its mean $\bar{R}$ and variance v_R, the case volume m and the balance scale defined below, are latent and refer to one generation of transmission. The weekly statistics computed from admissions in Section 3.2, and the intercept a, slope b and ratio b/a fitted to their relative loss, are apparent quantities defined on the scale of each statistic and on weeks. The ratio b/a would equal the balance scale only if a week were a single generation and the case-count scale were free of noise (Section 3.3). We therefore call b/a the apparent balance ratio, reserve $m_\times$ for the latent balance scale, and link the two through the calibration of Section 3.10.

Transmission potential varies between periods with climatic conditions, school terms and behaviour (Cauchemez et al., 2008; Shaman and Kohn, 2009; Dalziel et al., 2018). Its marginal moments are denoted $\mathbb{E}(R_t) = \bar{R}$ and $\mathrm{Var}(R_t) = v_R$. Because this variation acts on every chain of transmission in the same period, it corresponds to environmental stochasticity in population ecology, whereas the offspring and reporting variation of Eq. (2) corresponds to demographic stochasticity (Engen et al., 1998; Lande et al., 2003).

The scaled count $\hat{R}_t = C_t/m$ can be used for two different targets. The first is the current value R_t, which is the quantity of interest for situational awareness and for judging whether transmission is currently above one. The second is the mean across periods, $\bar{R}$, which is the quantity of interest when describing typical transmission in a season. Two related targets, the mean of the realised values over a finite set of periods and the value in the next period, are treated in Corollary 1. The error for the first target is $\hat{R}_t - R_t$. The error for the second target can be written as:

$$\hat{R}_t - \bar{R} = (\hat{R}_t - R_t) + (R_t - \bar{R}), \tag{3}$$

where the second term does not depend on m. Conditional on R_t, $\hat{R}_t$ is unbiased for R_t, and its variance $R_t/m + \rho {R_t}^2/(km)$ falls to zero as m grows; in this case $\hat{R}_t$ attains the Cramér–Rao bound for R_t (Appendix A). Because the

first term of Eq. (3) has mean zero given R_t, the two terms are uncorrelated, and the mean squared errors of the two targets differ by exactly v_R. This identity requires three conditions: the estimate is conditionally unbiased for R_t, the case volume m does not depend on R_t (H4), and ρ is fixed and known (H1). For the cross-period target, the true variation of R_t between periods is therefore part of the error; it is sampling variation of that target, not counting error, and it does not fall with m.

2.2. Error variance and the balance scale

For the second target, the error variance follows from the law of total variance applied to Eq. (3).

Proposition 1 (Error variance for the mean across periods)**.** *Under H1–H4, the mean squared error of $\hat{R}_t = C_t/(\rho n)$ about R_t is A/m, and its variance about $\bar{R}$ is given by:*

$$V(n,\rho) = v_R + \frac{\bar{R}}{\rho n} + \frac{\bar{R}^2 + v_R}{kn} = v_R + \frac{A}{m}, \qquad A = \bar{R} + \frac{\rho(\bar{R}^2 + v_R)}{k}, \tag{4}$$

where the three terms are the environmental, reporting and offspring contributions. Then:

(i) for fixed ρ, $V \to v_R$ as $n \to \infty$;

(ii) for fixed n, $V \to v_R + \bar{R}/n + (\bar{R}^2 + v_R)/(kn) > v_R$ as $\rho \to 1$;

(iii) at the balance scale $m_\times = A/v_R$, the sampling and environmental contributions are equal, and doubling m reduces V by the fraction $1/[2(1 + m/m_\times)]$;

(iv) among estimators that are conditionally unbiased for R_t, no estimator has a smaller mean squared error for $\bar{R}$ than V.

The balance scale plays the role of the critical population size in ecological models, at which demographic and environmental contributions to the variance of population growth are equal (Lande, 1993; Lande et al., 2003). It is a reference point for comparing the two contributions and does not by itself define a minimum operating size. Part (iv) is restricted to conditionally unbiased estimators. An estimator that shrinks $\hat{R}_t$ towards a prior mean can have smaller error for $\bar{R}$ (Appendix B.5).

Parts (i) and (ii) show that case volume can be raised in two ways with different limits. Coverage n removes both sampling terms in the limit, whereas full reporting leaves the offspring term. Near the baseline, a small proportional increase in n reduces V by $A/\bar{R}$ times as much as the same proportional increase in ρ.

The single-period result extends to estimates that combine several periods. Let $\bar{R}_T = T^{-1}\sum_{t=1}^{T} R_t$ denote the mean of the realised values over T periods, and let R_t follow either independent draws or a stationary first-order autoregressive process with autocorrelation ϕ.

Corollary 1 (Three targets over several periods)**.** *Under H1–H4, with counts independent across periods given (R_t), the mean $T^{-1}\sum_{t=1}^{T}\hat{R}_t$ and the one-step forecast of R_{t+1} have the following mean squared errors.*

(a) About the process mean $\bar{R}$, the error is $v_R/T + A/(Tm)$ for independent periods and $v_R S(\phi,T)/T^2 + A/(Tm)$ under the autoregressive process, where:

$$S(\phi,T) = \frac{T(1+\phi)}{1-\phi} - \frac{2\phi(1-\phi^T)}{(1-\phi)^2}. \tag{5}$$

(b) About the finite-period mean $\bar{R}_T$, the error is $A/(Tm)$ whatever the dependence of R_t.

(c) The naive forecast $\hat{R}_t$ of R_{t+1} has error $A/m + 2v_R$ for independent periods and $A/m + 2(1-\phi)v_R$ under the autoregressive process. With $\bar{R}$ known, the best linear forecast $\bar{R} + \lambda(\hat{R}_t - \bar{R})$, with $\lambda = \phi v_R/(v_R + A/m)$, has error:

$$v_R\left[1 - \phi^2\,\frac{m}{m + m_\times}\right], \tag{6}$$

which falls towards $(1-\phi^2)v_R$ as m grows and equals v_R for independent periods.

Part (b) shows that the finite-period mean behaves like the current value: its error is counting error alone and vanishes with case volume. Parts (a) and (c) show that targets beyond the observed periods keep an environmental floor. No linear forecast built from $\hat{R}_t$ has error below $(1-\phi^2)v_R$, and the naive forecast has error above $2(1-\phi)v_R$. The limit v_R also assumes that environmental variation is independent between places. If J jurisdictions share environmental variance v_R with mean pairwise correlation $\bar{c}$, the environmental variance of their average is given by:

$$v_R\,\bar{c} + \frac{v_R(1-\bar{c})}{J}, \tag{7}$$

where pooling reduces this limit only to the extent that variation is not synchronised ($\bar{c} < 1$).

2.3. Reporting heterogeneity and the allocation of a surveillance budget

Reporting completeness is rarely constant. Assumption H1 can be replaced by a reporting probability p_t that varies between periods independently of R_t and n, with mean ρ and relative volatility $\tau = \mathrm{SD}(p_t)/\rho$, and with counts independent across periods conditional on R_t and p_t (H5).

Proposition 2 (Reporting heterogeneity)**.** *Under H2–H5, the error variance of $\hat{R}_t = C_t/m$ about $\bar{R}$ can be written as:*

$$V_H(m) = v_R + v_H + \frac{A_H}{m}, \qquad v_H = \tau^2(\bar{R}^2 + v_R), \qquad A_H = \bar{R} + \frac{\rho(1+\tau^2)(\bar{R}^2 + v_R)}{k}, \tag{8}$$

where v_H is the variance added by reporting heterogeneity. The limiting variance rises to $v_R + v_H$, and the balance scale becomes $m_{\times,H} = A_H/(v_R + v_H) \leq m_\times$, with equality if and only if $\tau = 0$. With $Q = \bar{R}^2 + v_R$, the difference is given by:

$$m_\times - m_{\times,H} = \frac{\tau^2 Q\, \bar{R}\,(1 + \rho\bar{R}/k)}{v_R(v_R + v_H)}. \tag{9}$$

The variances v_R and v_H cannot be separated using the mean and variance of the counts alone, which identify only $(1+\tau^2)(\bar{R}^2 + v_R)$.

Proposition 2 is the identification limit for the empirical part of this study. An intercept estimated from counts alone combines transmission and reporting variation unless reporting is measured separately. The same difficulty arises in state-space models of population dynamics, where process and observation variance are often weakly identified from a single series (Knape, 2008; Auger-Méthé et al., 2016).

The balance scale also determines how a fixed budget is best divided. Suppose that T periods are observed with case volume m in each, that each period costs c_0 and each unit of case volume costs c_1, and that the budget is $\mathcal{B} = T(c_0 + c_1 m)$.

Proposition 3 (Allocation of a fixed budget)**.** *Under H1–H4 with R_t independent across periods, linear costs and T treated as continuous, the error of the mean of $\hat{R}_t$ about $\bar{R}$ over T periods is given by:*

$$\frac{v_R}{T} + \frac{A}{Tm}, \tag{10}$$

which, under the budget constraint, is minimised at:

$$m^* = \sqrt{\frac{c_0}{c_1}\, m_\times}, \tag{11}$$

where $T^* = \mathcal{B}/(c_0 + c_1 m^*)$, subject to $T^* \geq 1$. At a fixed budget the error increases with m if and only if $m > m^*$.

The optimal case volume per period is the geometric mean of c_0/c_1 and the balance scale, the same form as the optimal number of secondary units in two-stage sampling (Cochran, 1977) and the optimal cluster size in cluster-randomised trials (Raudenbush, 1997). Under the three conditions of the proposition, $m > m^*$ is necessary and sufficient for a reallocation of budget from cases to further periods to reduce error, because the error at fixed budget is a strictly convex function of m. The rule is not general. With integer T the optimum moves to a neighbouring feasible design, with non-linear costs it moves with the shape of the cost function, and with autocorrelated environments the environmental term becomes $v_R S(\phi, T)/T^2$ (Corollary 1) and the optimum depends on the budget. Appendix B gives the budget threshold at which the autoregressive optimum switches.

Reporting heterogeneity can reverse the ranking of coverage and completeness as targets for investment.

Corollary 2 (Coverage versus completeness)**.** *If τ does not depend on ρ, increasing coverage n reduces V_H more than the same proportional increase in ρ. If reporting variation arises from binomial sampling of ν units with $\mathrm{Var}(p_t) = \rho(1-\rho)/\nu$, increasing ρ is preferable whenever:*

$$m > \frac{\rho^2(\nu - 1)}{k}, \tag{12}$$

where ν is the number of reporting units sampled per period. Under a Beta prior $p_t \sim \mathrm{Beta}(\rho\nu, (1-\rho)\nu)$ with variance $\rho(1-\rho)/(\nu+1)$, the threshold in Eq. (12) becomes $m > \rho^2\nu/k$.

3. Materials and methods

3.1. Data and analysis windows

Weekly laboratory-confirmed influenza hospital admissions and the percentage of hospitals reporting them were taken from the NHSN Weekly Hospital Respiratory Data table for the 50 US states and the District of Columbia (Centers for

Disease Control and Prevention, 2026), using the version retrieved on 2026-10-20. The table contained 21641 jurisdiction-weeks from 2020-08-08 to 2026-10-10 and listed 67 geographic aggregations. Removing the ten Department of Health and Human Services regions and the national total left 56 jurisdictions, and removing American Samoa, Guam, the Northern Mariana Islands, Puerto Rico and the US Virgin Islands left the 50 states and the District of Columbia. Reporting of influenza admissions was required from 2 February 2022 to 30 April 2024 (Centers for Disease Control and Prevention, 2022; Frutos et al., 2026), voluntary from 1 May to 31 October 2024 and required again from 1 November 2024 (Centers for Medicare and Medicaid Services, 2024). The table gave 6009, 420 and 4077 jurisdiction-weeks with a case-count scale of at least five before, during and after the voluntary period; the first group includes weeks before the requirement began, and all analysis windows lie after it.

The primary windows covered early epidemic growth in four seasons from 2022–23 to 2025–26. Each window contained the contiguous weeks from 1 October, or from 9 November 2024 when mandatory reporting resumed, in which national weekly growth was at least 1.05, and ended the week before the first national turning point. The windows ran from 2022-10-01 to 2022-11-26, 2023-10-07 to 2023-12-23, 2024-11-09 to 2024-12-28 and 2025-10-04 to 2025-12-27, giving 42 calendar weeks. This rule replaced a first version that ended each window at the season maximum and gave only four weeks in the double-peaked 2024–25 season; the revision was made on 2026-10-20 before any intercept was fitted on the new windows, and no estimates under the first version were retained (Table S9). For continuity, the three calendar windows of an earlier version of this analysis (2022-10-08 to 2022-11-26, 2023-10-14 to 2023-12-16 and 2024-11-09 to 2025-01-18) were analysed as a secondary specification. A peak window contained the national peak week of each season and the seven weeks after it. No zero-imputed counts entered any window.

3.2. Weekly growth statistics

Two statistics of weekly growth were computed for each state and week. The weighted growth ratio is given by:

$$G_t = \frac{C_{t+1}}{D_t}, \qquad D_t = 0.65\,C_t + 0.25\,C_{t-1} + 0.10\,C_{t-2}, \tag{13}$$

where C_t is the number of admissions in week t and D_t is the case-count scale, with $\sum_s {w_s}^2 = 0.495$. Only weeks with $D_t \geq 5$ were used for G. The renewal-type ratio converts single-week growth into a reproduction number through the generation interval (Wallinga and Lipsitch, 2007), and is given by:

$$R_{A,t} = (1 + r_t\theta)^{\alpha}, \qquad r_t = \frac{1}{7}\log\frac{C_{t+1} + 0.5}{C_t + 0.5}, \tag{14}$$

where r_t is the daily growth rate, and α and θ are the shape and scale of a gamma distribution with mean 3.6 days and standard deviation 1.6 days. These values are the household serial interval of seasonal influenza estimated by Cowling et al. (2009), used here as a proxy for the generation interval. A serial interval is not a generation interval, and direct estimates of the influenza generation interval are shorter. A household study in the United States estimated an intrinsic mean of 3.2 days and a realised household mean of 2.8 days (Chan et al., 2025), and pooled serial intervals were 2.2 days for influenza A(H3N2) and 2.8 days for A(H1N1)pdm09 (Vink et al., 2014). The statistic was recomputed with intervals of mean 2.2 and 4.3 days as sensitivity analyses, and the consequences of a shorter true interval for the inversion were examined by simulation (Section 3.10). Because the scale of R_A is the single-week count, it was computed only for weeks with $C_t \geq 5$ as well as $D_t \geq 5$. R_A is a transformation of observed growth and not a posterior estimate of R_t (Cori et al., 2013); both statistics are therefore apparent quantities whose relationship to v_R was established by simulation (Section 3.10).

3.3. Expected loss and the balance scale

In routine surveillance the number of primary cases is unobserved, and secondary infections arise from the infection pressure $\Lambda_{t+1} = \sum_{s\geq 1} w_s I_{t+1-s}$ (Fraser, 2007; Cori et al., 2013). Applying the law of total variance to C_{t+1} given the history $\mathcal{F}_t$, with $m = \rho\mathbb{E}(\Lambda_{t+1} \mid \mathcal{F}_t)$, the relative prediction variance can be written as:

$$\frac{\mathrm{Var}(C_{t+1} \mid \mathcal{F}_t)}{m^2} = v_R + \frac{A}{m} + \mathrm{CV}^2(\Lambda_{t+1} \mid \mathcal{F}_t)\,(\bar{R}^2 + v_R), \tag{15}$$

where CV is the coefficient of variation. The last term is a further scale-invariant component that vanishes only when infection pressure is known. For a statistic X_t with mean $\bar{X}$ pooled over all state-weeks of a specification, the relative loss $L_t = [(X_t - \bar{X})/\bar{X}]^2$ was modelled as:

$$\mathbb{E}(L_t \mid D_t) = a + \frac{b}{D_t}, \tag{16}$$

where, if a week were a single generation and D_t were free of noise, $a = v_R/\bar{R}^2$, $b = A/\bar{R}^2$ and $b/a = m_\times$. For R_A, whose scale is the single-week count, C_t replaces D_t in Eq. (16). Because $\bar{X}$ is pooled, differences in mean growth

between seasons enter the relative loss and therefore the intercept; Table S5 gives season-specific intercepts computed with each season's own mean. The model also assumes one intercept for all states, so any dependence of the intercept on the size of a state is not separated from the slope (Section 5). Binomial noise in the denominator, with relative variance $(1-\rho)\sum_s {w_s}^2/m$, increases the slope to approximately:

$$b \approx \frac{A + (1-\rho)(\bar{R}^2 + v_R)\sum_s {w_s}^2}{\bar{R}^2}. \tag{17}$$

Theoretical ranges of b/a for G used $\bar{R}$ equal to the mean of G in each specification, $v_R \in \{0.04, 0.08\}$, $k \in \{0.2, 0.5\}$ (Lloyd-Smith et al., 2005) and an admission probability per infection $\rho \in \{0.0056, 0.021\}$, that is 178.6 and 47.6 infections per admission (Tokars et al., 2018), with and without Eq. (17). With $\bar{R} = 1.671$, the mean of G in the calendar windows, the uncorrected balance scale A/v_R over these scenarios was 21.3 to 49.2 cases per generation; with the rule-window mean of 1.724 it is 22.0 to 51.0. This range ignores both the conversion from weeks to generations and the calibration of Section 3.10, and serves only as a scenario range for the theoretical calculations.

3.4. Estimation and inference

Eq. (16) was fitted by feasible generalised least squares with weights $1/\hat{\mu}^2$, where $\hat{\mu}$ is the fitted mean from the previous iteration, iterated to convergence (Wooldridge, 2010). Observations from the same state are dependent over time, and observations from the same week share national influences. Uncertainty was therefore estimated by a two-way (pigeonhole) bootstrap that draws states and weeks independently with replacement and weights each observation by the product of the two draw counts (Owen, 2007). For crossed random effects no resampling scheme reproduces the variance exactly, and this bootstrap tends to overstate it slightly (McCullagh, 2000; Owen, 2007); multiway clustering motivates the same two-way structure in regression inference (Cameron et al., 2008, 2011). Each specification used 2000 bootstrap replicates with its own seed (20261024 for R_A and 20261124 for G in the rule windows, 20261224 for G and 20261324 for R_A in the calendar windows), and uncertainty is reported as the point estimate followed by the 95% percentile interval, for example 0.0150 (95% CI 0.0100 to 0.0208).

The intercept was decomposed by refitting Eq. (16) twice. The first refit removed national-week effects by using deviations from the mean across states in each week, with a small-sample factor $n_w/(n_w - 1)$. The second removed additive week and state-season effects estimated by alternating projections. The national-week component is the difference between the unadjusted and week-adjusted intercepts, and the state-season component is the further change on adding state-season effects. The affine model was compared with a power law $\mathbb{E}(L_t \mid D_t) = c{D_t}^{-\gamma}$ by a weighted contrast of residual sums of squares, calibrated by a residual bootstrap under the power law, and by leave-one-state-out and leave-one-season-out cross-validation. Power to detect a positive intercept was computed by simulating from the fitted affine model with intercepts up to the estimated value. Sensitivity analyses of b/a trimmed the most influential 1% of losses, omitted one season at a time, used Huber weights and removed national-week effects. Further analyses used the peak window, changed the generation interval and changed the denominator weights to a single week or three equal weeks.

3.5. Within-state fits

The intercept of Eq. (16) is common to all states, so the cross-sectional slope combines variation of the loss with case volume within states and differences between states. To separate the two, the loss of R_A in state s and week t was modelled as:

$$\mathbb{E}(L_{st} \mid C_{st}) = a_s + \frac{b_{\mathrm{w}}}{C_{st}}, \tag{18}$$

with 51 state intercepts a_s and a common within-state slope b_{w}, fitted by the same feasible generalised least squares with variance proportional to the squared fitted mean. The within-state ratio is $b_{\mathrm{w}}/\bar{a}_s$, where $\bar{a}_s$ is the mean of the state intercepts; the median, the mean weighted by the number of state-weeks and the mean weighted by the total count of each state were used as alternatives. The case-count scale D_t replaced C_t as a sensitivity analysis. Two further fits replaced the state intercepts by week intercepts and added week intercepts to the state intercepts. The dependence of a_s on the size of a state was summarised by the correlation of a_s with the logarithm of the state median of C_t. Uncertainty came from 1000 draws of the two-way bootstrap of Section 3.4. In each draw the cross-sectional and within-state models were both fitted, and the paired ratio of the cross-sectional ratio b/a to the within-state ratio was formed within the draw.

3.6. Binomial thinning

If the term of the loss that falls with case volume arises from Poisson-type counting, removing cases at random should raise the slope in a known way while leaving the intercept unchanged. On the fixed sample of 1748 state-weeks of the primary R_A specification, each weekly count was thinned once per replicate, $C^{(q)} \sim \mathrm{Bin}(C, q)$, for retention probabilities $q \in \{0.8, 0.6, 0.5, 0.4, 0.3, 0.2, 0.1\}$, with 400 replicates per value. Adjacent state-weeks shared the thinned count of their common week, and no window or filter was reapplied. In each replicate R_A was recomputed with the correction of 0.5

in Eq. (14), its pooled mean was re-estimated, and Eq. (16) was refitted with the original C_t as the scale. A delta-method expansion of Eq. (14) predicts that the fitted coefficients satisfy:

$$b(q) = b(1) + c\left(\frac{1}{q} - 1\right), \qquad a(q) = a(1), \tag{19}$$

where the Poisson value of c is b_P, the mean of ${\kappa_A}^2(1 + C_t/C_{t+1})$ over the sample, and κ_A is defined in Section 3.7. The rate c was estimated as the least-squares slope through the origin of $b(q) - b(1)$ on $1/q - 1$, together with the slope c_q at each level and the change in intercept $a(q) - a(1)$. Uncertainty came from 1000 draws of the two-way bootstrap, with one thinning draw per level in each bootstrap draw, so that changes in intercept were paired within draws. The analysis was repeated in the 828 state-weeks with $C_t \geq 50$, for which $qC_t \geq 5$ at every level. The mean loss over the sample was also fitted as $\alpha + \beta/q$, and $\alpha/(\alpha + \beta)$ gives the share of the unthinned mean loss that does not scale with $1/q$. Thinning changes only the reporting probability and leaves the transmission process, the timing of admissions and the composition of the sample unchanged. It therefore tests the scaling of the A/m term and does not identify v_R. Seeds were 20261102 for the thinning replicates, 20261103 for the bootstrap of the thinning analysis and 20261104 for the bootstrap of the within-state fits.

3.7. Reporting completeness

Coverage-adjusted statistics were computed after dividing each count by the share of hospitals reporting in that state and week. The change in intercept, $\Delta a = a - a_{\text{adj}}$, was tested for equivalence to zero by two one-sided tests at the 5% level, using 90% intervals from a paired two-way bootstrap (Schuirmann, 1987). The margins were set at the intercept that a relative reporting volatility τ_0 would add under Proposition 2, and are given by:

$$\delta_G = {\tau_0}^2\left(1 + \frac{v_R}{\bar{R}^2}\right), \qquad \delta_{R_A} = {\kappa_A}^2\,{\tau_0}^2\left(1 + \frac{v_R}{\bar{R}^2}\right), \tag{20}$$

where $\kappa_A = \mathrm{d}\log R_A/\mathrm{d}\log g$ propagates a perturbation of weekly growth g to R_A, $\bar{R}$ is the sample mean of each statistic, $v_R = 0.04$, and $\tau_0 = 0.10$ (primary) or 0.05. Because $\bar{R}$ is estimated, the margins expressed as shares of the intercept carry bootstrap intervals. The change expressed as a share of the intercept, $u = \Delta a/a$, was computed in each bootstrap replicate, and its one-sided upper bound is the 95th percentile of these values. Power was computed from a normal approximation. An upper bound u can be converted to a bound on reporting volatility in two ways. Against the equivalence margins, which scale with ${\tau_0}^2$, it corresponds to $\tau = \tau_0\sqrt{u/u_0}$, where u_0 is the margin expressed as a share of the intercept. On the latent scale, if reporting volatility accounted for the share u of the limit $v_R + v_H$, Proposition 2 gives $v_H = uv_R/(1 - u)$, and the corresponding volatility is given by:

$$\tau = \sqrt{\frac{u}{1-u}\,\frac{v_R}{\bar{R}^2 + v_R}}. \tag{21}$$

The margin-based conversion and Eq. (21) both refer to variation in facility participation only. The interquartile range of $\log G$ and the standard deviation of G over May to October were compared between 2023, 2024 and 2025. As a second proxy, R_A was computed from weekly positive influenza tests in clinical laboratories for 44 states with available data, and the share of the hospital intercept removed by regressing hospital R_A on laboratory R_A was estimated with 500 two-way bootstrap replicates.

3.8. Seasons and pathogens

Season-specific intercepts were pooled with a DerSimonian–Laird random-effects model (DerSimonian and Laird, 1986), with a t interval on $S - 1$ degrees of freedom for the pooled intercept and a Q-profile interval for the between-season variance (Viechtbauer, 2007). Corollary 3 gives the lower bound on the precision of any between-season variance (Searle et al., 1992). Intercept ratios $q_1 = a_{(1,0,0)}/a_w$ and $q_2 = a_{(1/3,1/3,1/3)}/a_w$ compare a single-week and an equal three-week denominator with the weights of Eq. (13). They were computed for influenza in four windows and for COVID-19 and RSV admissions in two windows each. Null ranges came from simulated epidemic growth with reporting noise but no environmental variation, over three generation intervals, three admission delays and four noise levels. Each of the eight combinations was tested on five one-sided criteria ($q_1 < 1$, $q_2 > 1$, each ratio outside its null range and a positive excess over $v_R/\bar{R}^2$), and the 40 tests were adjusted by the Holm and Benjamini–Hochberg procedures (Holm, 1979; Benjamini and Hochberg, 1995).

3.9. Simulation of apparent statistics

The closed forms of Section 2 were checked by Monte Carlo with 20000 replications per point. The difference of v_R between the two mean squared errors was checked with 200000 draws at each of 22 case volumes between 2 and 5000 under the baseline of Table 1, and Corollary 1 at $\phi \in \{0, 0.5, 0.8\}$, $m \in \{10, m_\times, 2000\}$ and $T = 10$. Proposition 2 was checked with beta-distributed reporting at $\tau = 0.15$, and the optimum of Proposition 3 by a grid search with continuous

T. Under first-order autocorrelation the budget-optimal case volume was found numerically for $\phi \in \{0.5, 0.8\}$, $c_0/c_1 \in \{10, 100\}$ and budgets from 300 to 30000; these checks used the seed 20261107. In a first, generic design, apparent weekly statistics were simulated from a daily renewal model with $\bar{R} = 1.2$ and the generation interval above. Weekly R_t was drawn from a gamma distribution with variance v_R, independently across weeks, offspring were negative binomial and admissions were binomial with probability ρ_h after a gamma delay with mean 5.5 days and standard deviation 3.0 days, which was assumed. Each replicate contained 500 states with initial case-count scales log-uniform between 2 and 150, four warm-up weeks and six observation weeks. Seventeen scenarios crossed $v_R \in \{0.01, 0.02, 0.04, 0.08, 0.16\}$ with $k \in \{0.1, 0.35, 1\}$ at $\rho_h = 0.01$ and added $\rho_h \in \{0.005, 0.02\}$ at the baseline, each with 20 replicates. The noise-free limit of each intercept came from a mean-field run of 20000 states in which all count noise was replaced by expectations.

3.10. Two-step inversion

The empirical intercepts were mapped to v_R in two steps. In the first step, a data-matched version of the simulator gave the expected fitted intercept of each statistic as a function of v_R. Each simulated panel had the cross-section and length of the analysis panel, with 51 states and the four rule windows of 9, 12, 8 and 13 observation weeks; the window rule was not reapplied to simulated national growth. The season mean R_s was the renewal-type transform of the geometric-mean national weekly growth over each window, and initial admissions in each state equalled its observed admissions in the first window week. Weekly reproduction numbers were $R_{i,s,w} = R_s + \eta_{s,w} + \varepsilon_{i,s,w}$, where the national week shock η has variance $0.45\, v_R$ and is shared by all states, the state shock ε has variance $0.55\, v_R$, and values were truncated below at 0.05. The share of 0.45 was set between the national-week shares of 0.41 and 0.51 observed for the two statistics (Section 4.7); matching the simulated intraclass correlation and national-week share to the data is therefore a consistency check and not an independent validation, and setting the share to zero or 0.65 changed the implied variance for R_A from 0.040 to 0.039 and 0.040 (Table S8). Transmission was a daily renewal process with negative binomial offspring ($k = 0.35$), admissions were binomial with probability $\rho_h = 0.01$ after the delay above, and both statistics were computed and filtered exactly as in the data, with the analyst's generation interval of mean 3.6 days. The design crossed $v_R \in \{0.02, 0.04, 0.06, 0.09\}$ with 200 panel replicates per value, and Eq. (16) was fitted to each replicate. A calibration line $\mathbb{E}(\hat{a}) = \gamma_0 + \gamma_1 v_R$ was fitted by weighted least squares with weights equal to the inverse Monte Carlo variance of each mean, and the scale factor $\kappa = \mathbb{E}(\hat{a})/(v_R/\bar{R}^2)$ was computed with $\bar{R}$ the row-weighted true mean.

In the second step, the implied variance $v_R^{\text{impl}} = (a - \gamma_0)/\gamma_1$ was computed for each statistic. Its interval combined 2000 joint two-way bootstrap draws of the two empirical intercepts, which share state and week multiplicities, with independent draws of (γ_0, γ_1) from their Monte Carlo sampling distribution. The interval is conditional on the simulator, including the generation interval, and does not include uncertainty about it. Agreement between the two statistics was defined as non-rejection, that is, a 95% interval of their difference that contains zero. For the primary design, equivalence was also tested by two one-sided tests on the ratio of the implied variances with bounds of 0.80 and 1.25, using the 90% interval of the ratio draws. Sensitivity settings changed one element of the design at a time (Table S8). They removed the national shock or raised its share to 0.65, added first-order autocorrelation of 0.5 across weeks to both the national and the state shocks, or set the true generation interval to a mean of 2.2 or 4.3 days while the analyst kept 3.6 days. Two settings changed the delay mean to 3.5 or 8 days. One further setting lowered k to 0.1, another replaced the transmission shock by a synchronous multiplicative shock to reported counts, and the generic design above served as a bridge. Four further settings are reported in Table S10. One set the true generation interval to 2.9 days, the lower end of the 95% confidence interval for the mean serial interval of Cowling et al. (2009) and close to direct estimates of the generation interval (Chan et al., 2025). Another set the first-order autocorrelation to 0.25, and two drew weekly admissions from a gamma–Poisson mixture with a variance-to-mean ratio of 1.5 or 2.0. The independence of weekly shocks was checked at the state level with the lag-1 autocorrelation of $x/\bar{x} - 1$ pooled within state-season blocks of each statistic, with an interval from a state-cluster bootstrap of 1000 draws, and compared with its distribution over simulated panels at $v_R = 0.04$. This check does not test autocorrelation or trend in the series of national-week means, and no setting gave autocorrelation to the national shock alone. The calibrated theoretical b/a used the same design over $v_R \in \{0.04, 0.08\}$, $k \in \{0.2, 0.5\}$ and $\rho_h \in \{0.0056, 0.021\}$, and the range of b/a across single replicates gave the distribution expected from one panel of this size. Simulation seeds were derived from 20261020 by a seed sequence indexed by setting, value of v_R and replicate block; the joint bootstrap and the calibration draws used 20261921 and 20261922.

4. Results

The results are arranged in three layers. Sections 4.1 and 4.2 check the theory in simulations where the truth is known. Sections 4.3 to 4.5 describe how the loss of a weekly statistic varies with case volume in NHSN admissions and test that relation by thinning the counts and by fitting state intercepts. Sections 4.6 to 4.9 calibrate the empirical intercepts against the latent model and trace the consequences for reporting and design.

4.1. Two targets with known truth

Under the baseline parameters ($\bar{R} = 1.15$, $v_R = 0.04$, $k = 0.35$, $\rho = 0.25$), the sampling coefficient was $A = 2.1232$ and the balance scale was $m_\times = 53.1$ reported cases per period (Fig. 1B). The mean squared error about the current

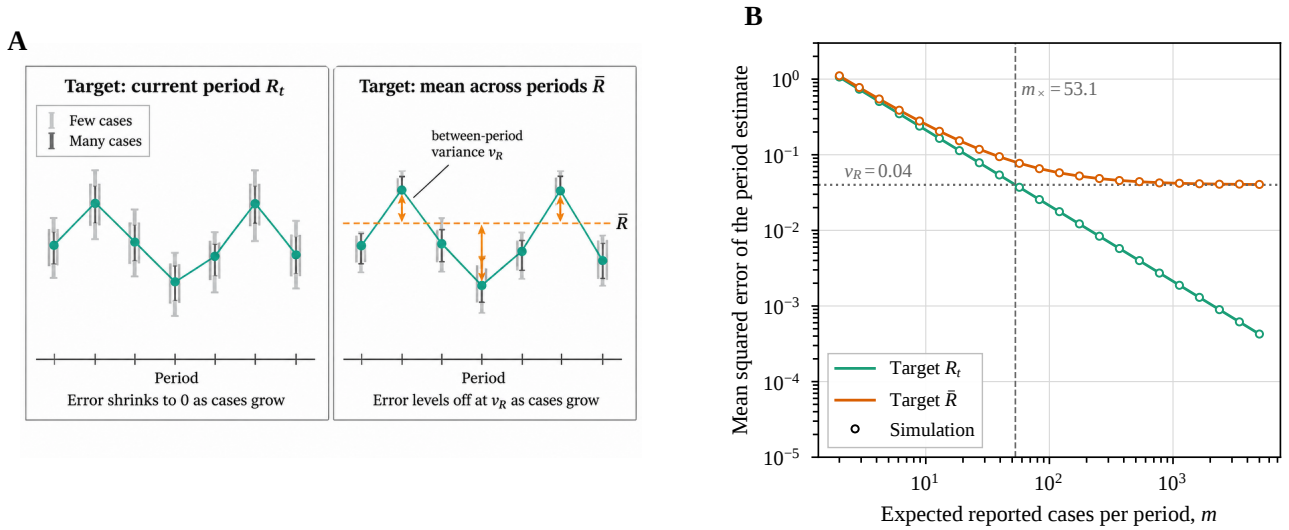


Fig. 1. Two inferential targets for a case-count estimate of the reproduction number. (A) Schematic. For the current value R_t, the interval of each period's estimate narrows as more cases are counted. For the mean across periods, $\bar{R}$, the deviations $R_t - \bar{R}$ (arrows), whose variance is v_R, remain however many cases are counted. (B) Mean squared error of $\hat{R}_t = C_t/m$ about R_t and about $\bar{R}$ against expected reported cases per period, at $\bar{R} = 1.15$, $v_R = 0.04$, $k = 0.35$ and $\rho = 0.25$. Lines are the closed forms A/m and $v_R + A/m$; circles are Monte Carlo estimates from 200000 draws per value, with 95% Monte Carlo intervals mostly hidden by the symbols. The dotted line marks v_R and the dashed line the balance scale $m_\times = 53.1$.

value fell in proportion to $1/m$ over the whole range, whereas the error about the mean levelled off at v_R. Across 22 case volumes from 2 to 5000, Monte Carlo estimates of both errors differed from A/m and $v_R + A/m$ by at most 0.74%, so the two curves were separated by v_R throughout. Doubling the case volume reduced the error about the mean by 42.1% at $m = 10$, by 25.0% at $m = m_\times$ and by 10.5% at $m = 200$.

The formulas of Corollary 1 for the process mean over ten periods, the finite-period mean and the naive and best linear one-step forecasts matched simulation to within 1.39% at $\phi = 0$, 0.5 and 0.8. Applying the formula for independent periods, $v_R/T + A/(Tm)$, to autocorrelated periods departed from the simulated error by up to 430%, which is the size of the correction $S(\phi, T)$ in Eq. (5).

Raising reporting probability at fixed coverage did not reduce error about the mean to v_R. At full reporting the residual variance was 0.166 with 40 primary cases, 0.0652 with 200 and 0.0450 with 1000, against $v_R = 0.04$. Near the baseline, a proportional increase in coverage reduced error 1.85 times as much as the same proportional increase in reporting probability.

4.2. Reporting heterogeneity and the allocation of a budget

A relative reporting volatility of $\tau = 0.15$ added $v_H = 0.0307$ to the limiting variance and reduced the balance scale from 53.1 to 30.4 cases (Fig. 2A). Eq. (9) reproduced this difference to rounding error over the range of τ examined, and simulated $V_H(m)$ with beta-distributed reporting differed from Eq. (8) by at most 0.23%.

With independent periods, linear costs and continuous T, a grid search over designs located the optimum within 0.03% of Eq. (11), and at every budget examined the error increased with m exactly when $m > m^*$. Under first-order autocorrelation the rule no longer held. For $\phi = 0.5$ and 0.8, cost ratios of 10 and 100 and budgets from 300 to 30000, the numerical optimum ranged from 0.33 to 2.63 times m^*, and it depended on the budget (Fig. 2B). The scenario parameters for influenza admissions give uncorrected balance scales of 21.3 to 49.2 cases per generation (Section 3.3); over this range, m^* was 14.6 to 22.2 cases per period for $c_0/c_1 = 10$ and 46.2 to 70.1 cases for $c_0/c_1 = 100$ (Fig. S1C).

In a hypothetical example with ten periods of 50 cases, autocorrelated environment and reporting ($\phi = \psi = 0.5$) and $\tau = 0.15$, the mean squared error of the mean was 22.57×10^{-3}, of which 10.40×10^{-3} was environmental, 7.88×10^{-3} reporting and 4.29×10^{-3} sampling (Table S7). A tenfold increase in case volume reduced it to 18.71×10^{-3}, whereas standardising reporting reduced it to 14.65×10^{-3} and four times as many periods to 6.17×10^{-3}. At twice the baseline budget, spending the increment on cases gave 19.35×10^{-3} and spending it on periods gave 11.98×10^{-3} (Fig. S1D).

4.3. Loss and case volume in NHSN admissions

In the rule-defined early-growth windows of four seasons, 1748 state-weeks in 42 calendar weeks gave a renewal-type ratio intercept of $a = 0.0150$ (95% CI 0.0100 to 0.0208), slope $b = 0.560$ (95% CI 0.367 to 0.769) and $b/a = 37.4$ (95% CI 19.1 to 71.0) (Fig. 3A; Table S1). The weighted growth ratio in the same windows gave $a = 0.0840$ (95% CI 0.0498 to 0.1248) and $b/a = 23.5$ (95% CI 8.4 to 59.7) from 1756 state-weeks. The eight additional state-weeks have a case-count scale D_t of at least five but fewer than five admissions in the current week, which R_A requires because its scale is the single-week count. In the three calendar windows of the earlier analysis, the growth ratio gave $a = 0.0991$

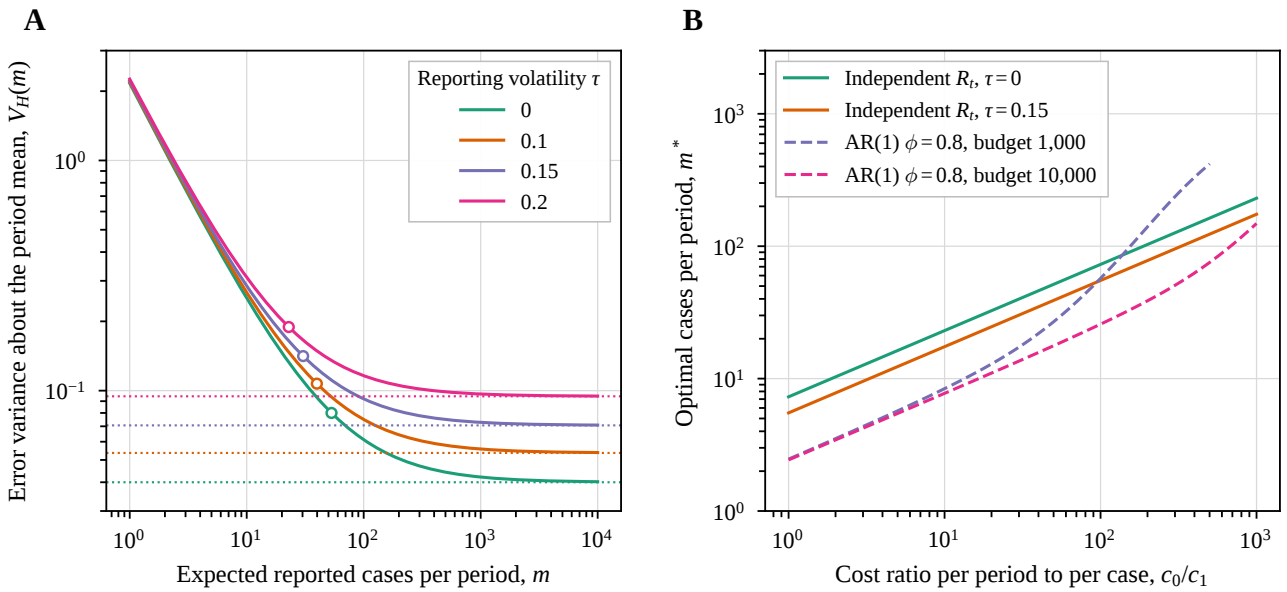


Fig. 2. Reporting heterogeneity and the allocation of a fixed budget. (A) Error variance about the mean across periods, $V_H(m)$, for relative reporting volatility τ of 0, 0.1, 0.15 and 0.2. Dotted lines mark the limits $v_R + v_H$ and open circles the generalised balance scale $m_{\times,H}$. (B) Optimal reported cases per period against the ratio of cost per period to cost per case, c_0/c_1. Solid lines give $m^* = \sqrt{(c_0/c_1)m_{\times,H}}$ for independent R_t at $\tau = 0$ and 0.15. Dashed lines give the numerical optimum under a first-order autoregressive R_t with $\phi = 0.8$ for budgets of 1000 and 10000 in units of the cost per case; designs with a single period are omitted.

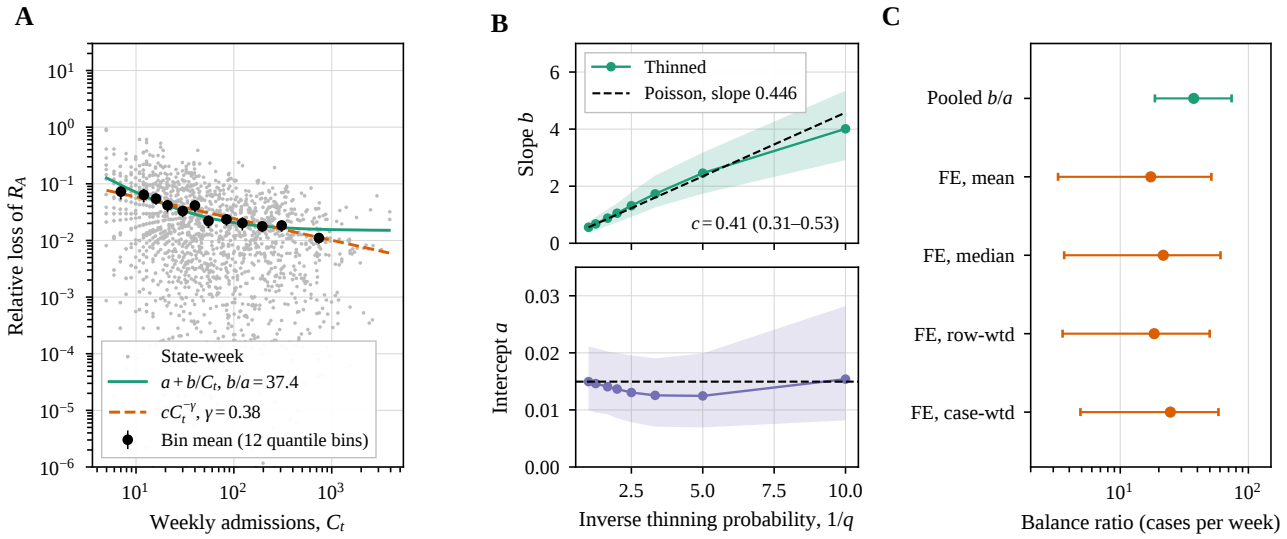


Fig. 3. Relation between loss and case volume in NHSN influenza admissions and two tests of it. (A) Relative loss of the renewal-type ratio R_A against weekly admissions C_t for 1748 state-weeks. Black points are means of 12 quantile bins with ± 1.96 standard errors that ignore clustering. The solid line is the affine fit $a + b/C_t$ ($b/a = 37.4$) and the dashed line the power law $cC_t^{-\gamma}$ ($\gamma = 0.38$). (B) Binomial thinning of the observed counts. Thinning changes only the reporting probability and leaves transmission, reporting volatility and the sample composition unchanged, so it tests the scaling of the term that falls with case volume and does not identify v_R. The upper panel gives the slope on the original scale, $b(q)$, against the inverse retention probability $1/q$, with the Poisson prediction $b(1) + 0.446(1/q - 1)$ as a dashed line. The lower panel gives the intercept $a(q)$, with $a(1)$ as a dashed line. Bands are 95% two-way bootstrap intervals. (C) Apparent balance ratio of the pooled cross-sectional fit and of fits with state intercepts (FE), in which the within-state slope is divided by the mean, median, row-weighted or case-weighted state intercept, with 95% two-way bootstrap intervals from 1000 draws.

(95% CI 0.0569 to 0.1382) and $b/a = 22.6$ (95% CI 9.0 to 52.2). Within the affine model, all bootstrap replicates gave a positive intercept for both rule-window specifications, as did 99.95% for the calendar windows. The bootstrap standard deviation of a was 1.95 to 4.68 times that obtained by resampling states only.

The ratio b/a varied across specifications by a factor of 3.9 for R_A, from 24.7 after trimming the most influential 1% of losses to 97.1 in peak weeks, and by a factor of 7.1 for G, from 12.2 with equal denominator weights to 86.8 in peak weeks (Fig. S3D). Generation interval variants changed it by less than 3% for R_A. Over the observed range of scales, a power law without an intercept, with exponent $\gamma = 0.38$ for R_A, fitted the binned losses about as well as the affine form (Fig. 3A; bootstrap $p = 0.114$ for R_A, 0.438 and 0.418 for G; Table S2). Cross-validation by state and by season did not

consistently favour the affine form, and a power analysis found no fraction of the estimated intercept detectable with 80% power (Table S2). The data therefore do not distinguish a positive floor from a loss that keeps declining as a power of the case-count scale. The affine form was retained as a working point because its parameters correspond to Proposition 1 and enter the calibration directly. Under the power law no balance scale exists, so the intercept, the apparent balance ratio and the quantities derived from them below are conditional on the affine form.

4.4. Binomial thinning

Thinning the weekly counts raised the slope on the original scale from 0.560 without thinning to 4.01 at $q = 0.1$ (Fig. 3B; Table S11). The pooled rate was $c = 0.407$ (95% CI 0.311 to 0.531), compatible with the Poisson value $b_P = 0.446$; the ratio c/b_P was 0.91 (95% CI 0.69 to 1.19). With a free intercept the rate was 0.386 (95% CI 0.267 to 0.530). The rate at each level was 0.47 to 0.50 for $q \geq 0.2$ and 0.383 at $q = 0.1$, where many thinned counts are small. In the 828 state-weeks with $C_t \geq 50$ it was 0.477 (95% CI 0.286 to 0.769), against a Poisson value of 0.420 in that subset.

No change in the intercept was detected. At every level the 95% interval of $a(q) - a(1)$ included zero, for example -1.3×10^{-3} (95% CI -4.8×10^{-3} to 1.8×10^{-3}) at $q = 0.5$ and 0.4×10^{-3} (95% CI -7.0×10^{-3} to 11.9×10^{-3}) at $q = 0.1$. These intervals are wide relative to $a(1) = 0.0150$, so the analysis excludes only large changes. Fitted as $\alpha + \beta/q$, the mean loss over the sample had $\alpha = 0.0301$ (95% CI 0.0189 to 0.0417) and $\beta = 0.0117$ (95% CI 0.0086 to 0.0154), so 72% (95% CI 58 to 81) of the unthinned mean loss did not scale with $1/q$.

Because $b(1) = 0.560$ exceeds c, the slope on the thinned scale, $qb(q)$, fell from 0.56 to 0.40 as q decreased. Under Proposition 1 thinning amplifies only the reporting part of A/m, whereas the offspring part, $(\bar{R}^2 + v_R)/(kn)$, does not change with ρ; the excess of $b(1)$ over c is therefore consistent with a contribution of offspring overdispersion to the slope. Thinning tests the scaling of the term that falls with case volume. It leaves transmission and reporting volatility unchanged and therefore cannot separate them within the intercept.

4.5. Fits with state intercepts

With 51 state intercepts the within-state slope was $b_{\mathrm{w}} = 0.379$ (95% CI 0.104 to 0.669) and the mean state intercept $\bar{a}_s = 0.0219$ (95% CI 0.0124 to 0.0351), giving a within-state ratio of 17.3 (95% CI 3.3 to 51.4), against 37.4 for the pooled cross-section (Fig. 3C; Table S11). The ratio was 21.7 with the median state intercept, 18.4 with the row-weighted mean and 24.6 with the case-weighted mean, and 17.7 (95% CI 3.2 to 51.4) with D_t as the scale. No state intercept was negative. Within the same bootstrap draws, the ratio of the cross-sectional to the within-state value was 2.16 (95% CI 0.93 to 7.75), so the difference is not distinguished from sampling variation.

State intercepts correlated negatively with the logarithm of the state median count, $r = -0.44$ (95% CI -0.67 to 0.08). Larger states thus tended to have lower intercepts, as spatial averaging of environmental variation over more local epidemics would produce (Eq. (7)), but the interval includes zero. With week intercepts in place of state intercepts the slope was 0.508 (95% CI 0.232 to 0.835). With both sets of intercepts the slope was 0.78 (95% CI -1.20 to 2.77) and was not identified, because removing state and week means left 41% of the variance of $1/C_t$, against 74% after removing state means alone.

4.6. Calibration and the implied environmental variance

The empirical statistics are apparent quantities. The weighted growth ratio G compounds more than one generation of growth into a week, and the renewal-type ratio R_A maps weekly growth back to the generation scale through an assumed generation interval (Section 3.2). Their intercepts therefore need a calibration before they can be compared with v_R. We built this calibration in a simulation that reproduces the observed panel (Section 3.10). It keeps the four seasons, the rule-window lengths and the initial admissions of each state, and adds a national week shock that carries 45% of v_R and the delay from infection to admission. We simulated 200 panels at each of four values of v_R. The mean fitted intercept rose linearly with v_R for both statistics (Fig. S2A), and no cell mean departed from the fitted line by more than 0.60 Monte Carlo standard errors.

Expressed as a scale factor on $v_R/\bar{R}^2$, the fitted intercept was $\kappa = 0.516$ for R_A and 2.729 for G, close to the noise-free limits of 0.493 and 2.745 (Fig. S2B). A simpler design with 500 independent states and six observation weeks gave noise-free limits of 0.485 and 2.698, in line with the medians of 0.484 and 2.68 over the 17 scenarios of the first simulation. The data-matched panels also reproduced the dependence structure of the data. The intraclass correlation of weekly losses within calendar weeks was 0.258 for R_A and 0.310 for G in simulation, against 0.253 and 0.310 in the data, and the national-week share of the simulated R_A intercept was 0.43, against 0.41 observed. Across settings that changed one ingredient at a time, κ for R_A moved most with the true generation interval (1.287 for a mean of 2.2 days and 0.380 for 4.3 days, with the statistic still computed at 3.6 days), with autocorrelation of transmission between weeks (0.780 at $\phi = 0.5$) and with a synchronous shock to reported counts (1.041). It changed little with the national share of the shock, the admission delay or the offspring dispersion (0.454 to 0.535; Table S8).

Mapping each empirical intercept through its calibration line gave an implied latent variance of $v_R^{\mathrm{impl}} = 0.040$ (95% CI 0.026 to 0.056) for R_A and 0.043 (95% CI 0.025 to 0.064) for G in the baseline setting, in which the true mean generation interval equals the 3.6 days assumed by the statistic (Fig. S2A, C). Bootstrap variation in the empirical intercept accounted

for nearly all of this width, since the interval from calibration error alone was 0.0397 to 0.0406 for R_A. The scale factors of the two statistics differ by a factor of 5.3. Their implied variances differed by -0.003 (95% CI -0.012 to 0.005), a ratio of 0.93 (95% CI 0.78 to 1.18). The 90% interval of this ratio, 0.81 to 1.13, lies within 0.80 to 1.25, so two one-sided tests at the 5% level declare the two estimates equivalent within a margin of 20%. This agreement is a weak check on the calibration, because the two statistics also agreed in nine of the ten simulation settings, whose implied variances differed by a factor of up to 4.8 (Table S8). On the latent scale, Proposition 1 at the simulator values ($\bar{R} = 1.18$, $\rho = \rho_h = 0.01$, $k = 0.35$) gives a balance scale of 30.3 admissions per generation (95% CI 21.7 to 46.3) for R_A, or 58.9 (95% CI 42.3 to 90.0) per week at 1.94 generations per week. The apparent ratio of 37.4 is defined on the scale of R_A and does not estimate this latent quantity directly.

The implied variance depended most on the true generation interval, which the interval of the baseline estimate does not include (Fig. 4B; Table S8). For R_A it was 0.013 (95% CI 0.008 to 0.019) with a true mean of 2.2 days, 0.026 (95% CI 0.017 to 0.036) with 2.9 days and 0.058 (95% CI 0.038 to 0.081) with 4.3 days. With the baseline interval and first-order autocorrelation of 0.25 or 0.5 between weeks it was 0.033 (95% CI 0.021 to 0.047) and 0.027 (95% CI 0.017 to 0.040). Values for G were close to those for R_A in each setting, 0.014 and 0.027 at 2.2 and 2.9 days. Each setting changed one element of the simulator, so this envelope does not cover joint changes. The latent weekly balance scale moves inversely with the implied variance, from 35.3 admissions per week at 4.3 days to 110.8 at 2.9 days and 273.4 at 2.2 days. The 95% interval of the difference between the two statistics contained zero in nine of the ten settings of Table S8; the exception replaced the transmission shock by a synchronous multiplicative shock to reported counts, which gave 0.019 for R_A and 0.029 for G (ratio 0.67; 95% CI 0.39 to 0.90).

Three further analyses examined the assumptions of the inversion (Table S10). The lag-1 autocorrelation of deviations from state-season means was -0.095 (95% CI -0.168 to -0.019) for R_A and 0.123 (95% CI 0.039 to 0.201) for G. Both statistics are autocorrelated by construction even when weekly shocks are independent; simulated panels with independent shocks gave means of -0.074 and 0.114, which placed the data at the 31st and 62nd percentiles. With a first-order coefficient of 0.25 the data fell at the 3rd and 10th percentiles, and with 0.5 below all simulated panels. The mean household serial interval of seasonal influenza has been estimated at 3.6 days (95% CI 2.9 to 4.3) (Cowling et al., 2009), and direct estimates of the mean generation interval lie near the lower end of this interval (Chan et al., 2025). The autocorrelation of the data sat at the 2.5th and 9th percentiles of panels simulated with 2.9 days and at the 80th and 92nd percentiles with 4.3 days. This state-level check therefore does not favour the shorter interval, although it leaves the series of national-week means untested (Section 3.10). The shorter interval did reproduce the empirical b/a. At $v_R = 0.02$, the simulated value nearest its implied variance, the data sat at the 52.5th and 58th percentiles of the simulated b/a. Overdispersed weekly admissions with a variance-to-mean ratio of 1.5 or 2.0 kept the implied variance between 0.039 and 0.045. They raised the simulated b/a at $v_R = 0.04$ from 28.0 to 68.0 and 81.0 for R_A, and from 16.2 to 40.2 and 48.4 for G. These settings also gave lag-1 autocorrelations below the data for both statistics, so the data favour overdispersion smaller than a variance-to-mean ratio of 1.5.

On the scale of each statistic, the calibrated theory gave b/a of 13.9 to 31.1 for R_A and 8.0 to 18.8 for G over eight combinations of $v_R \in \{0.04, 0.08\}$, $k \in \{0.2, 0.5\}$ and admission probability $\rho_h \in \{0.0056, 0.021\}$ (Fig. S2D). The empirical point estimates of 37.4 and 23.5 lie above these ranges, and their 95% intervals overlap them. Ratios from single simulated panels varied widely; 95% of them lay between 8.8 and 42.7 for R_A and between 3.6 and 28.6 for G, and the empirical values fall at the 93rd and 91st percentiles of these distributions. Overdispersed admissions and a shorter generation interval can each raise the simulated ratio to the empirical values (Table S10). The simulator gives every state the same intercept, and the within-state ratio of 17.3 lies inside the calibrated range for R_A (Section 4.5); spatial averaging in larger states is therefore a third candidate explanation of the high cross-sectional ratio.

4.7. Multiscale decomposition of surveillance uncertainty

The intercept contains a component shared by all states in a calendar week and a component specific to states (Fig. 4A; Fig. S3C; Table S2). Removing national-week effects reduced a by 41% (95% CI 12 to 59) for R_A and by 51% (95% CI 21 to 67) for G in rule windows, and by 52% for G in calendar windows. These are shares of the intercept, the part of the loss that does not fall with case volume, and not of the total variance of the statistic. Removing state-season effects as well changed the intercept little; for R_A it was 0.0083 (95% CI 0.0064 to 0.0141) after both adjustments, and the interval for the state-season component included zero in every specification. The intraclass correlation of the loss within calendar weeks, which refers to the total loss, was 0.253 for R_A and 0.310 for G. The bootstrap reflected the same structure. Resampling weeks alone gave a standard deviation of $\hat{a}$ of 0.0021 for R_A and 0.0158 for G, 1.46 and 1.87 times the values from resampling states alone, and two-way resampling gave 0.0028 and 0.0193.

The national-week component enters the loss of every state in the same week. More admissions in one state reduce only that state's count term, and averaging over states reduces only the state-specific term, so neither removes the shared component; only further weeks average it out. Pooling the mean weekly deviation over J randomly chosen states reduced its variance to 0.39, 0.31 and 0.25 of the single-state value for R_A at $J = 5$, 10 and 51, and to 0.43, 0.36 and 0.31 for G (Fig. S5C). These values approach the national-week component, which is 0.23 of single-state variance for R_A and 0.29 for G, and come within 10% of it at $J = 33$ and $J = 25$ respectively, as Eq. (7) predicts for partly synchronised variation.

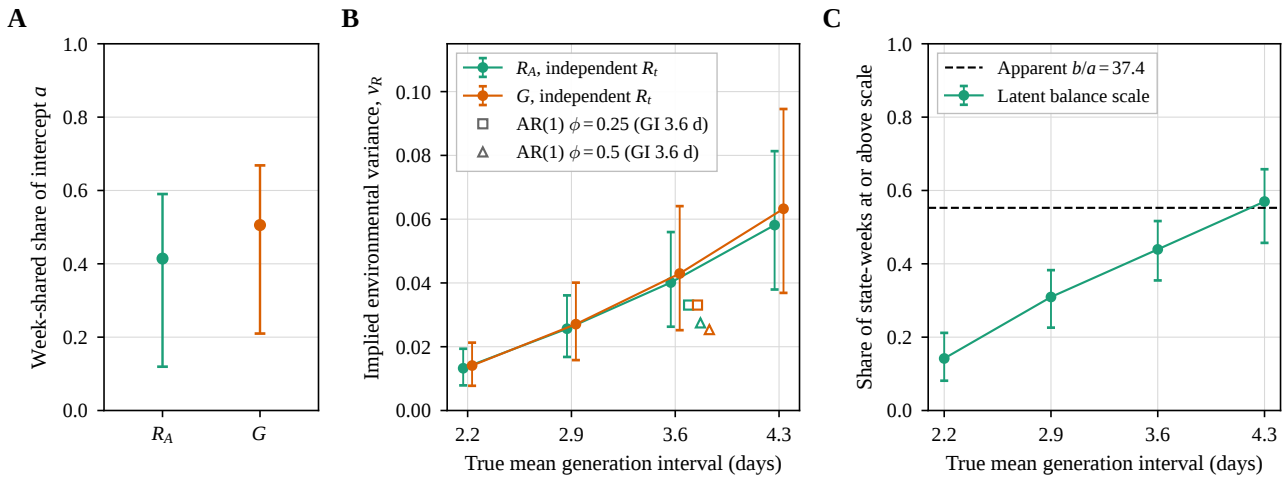


Fig. 4. Shared variance and the dependence of the implied environmental variance on the generation interval. (A) Share of the intercept a removed by national-week effects for R_A and G, with 95% two-way bootstrap intervals. (B) Implied environmental variance v_R against the true mean generation interval of the calibration simulation, with both statistics computed at a mean of 3.6 days. Filled points assume independent weekly R_t and carry 95% intervals; open symbols give first-order autocorrelation of 0.25 and 0.5 at 3.6 days, offset horizontally for legibility. (C) Share of the 1748 state-weeks of R_A with C_t at or above the latent weekly balance scale implied at each generation interval; intervals are the shares at the ends of the 95% interval of the scale. The dashed line is the share above the apparent ratio $b/a = 37.4$.

4.8. Mandatory and voluntary reporting regimes

Among state-weeks with a case-count scale of at least five, the median share of hospitals reporting was 92.3% in the first mandatory period, 45.7% under voluntary reporting and 92.9% after mandatory reporting resumed (Fig. S4A). In the rule windows the median was 93.7%, and 79.8% of state-weeks had at least 90% of hospitals reporting. Within these windows, adjusting each statistic for the reporting share changed the intercept of R_A by $\Delta a = 0.29 \times 10^{-3}$ (90% CI -0.22×10^{-3} to 0.85×10^{-3}) and that of G by 0.37×10^{-3} (90% CI -2.83×10^{-3} to 2.81×10^{-3}) (Fig. S4B; Table S3). Both changes were equivalent to zero within margins set at a relative reporting volatility of $\tau_0 = 0.10$. These margins equal 17% (95% CI 12 to 25) of the intercept of R_A and 12% (95% CI 8 to 20) of that of G, and they depend on the sample mean of each statistic through Eq. (20). The measured change was 2.0% of the intercept for R_A (95% CI -2.4 to 6.1; one-sided 95% upper bound 5.3%, the 95th percentile of the bootstrap distribution of the share) and 0.4% for G (95% CI -4.4 to 4.0; upper bound 3.5%), and all four seasons for R_A and three of four for G were equivalent individually. At $\tau_0 = 0.05$, with margins of about 4% and 3% of the intercepts, neither change was equivalent, and the power of the test at a true change of zero was 0.25 for R_A and 0 for G (Fig. S4C). Because most state-weeks in the windows had nearly complete participation, these tests bound the effect of small departures from full participation and do not extend to low participation.

The voluntary period gave the only contrast in which reporting completeness was low. Over May to October, the interquartile range of log G was 0.759 in 2024 under voluntary reporting, against 0.479 in 2023 and 0.483 in 2025 (Fig. S4D; Table S4). The ratio of 2024 to 2023 was 1.58 (95% CI 0.95 to 2.23) and that of 2024 to 2025 was 1.57 (95% CI 0.95 to 2.39), whereas the placebo ratio of 2025 to 2023 was 1.01 (95% CI 0.72 to 1.25). Median case-count scales in these months were similar (15.6, 16.6 and 16.2 admissions), and matching state-weeks on case-count scale gave ratios of 1.68 and 1.58. Matching on calendar month gave a median ratio of 1.33, with the ratio above one in nine of twelve month-by-year comparisons. COVID-19 admissions showed ratios of 1.62 against 2023 and 1.18 against 2025, and RSV admissions a ratio of 1.36 against 2025. All of these comparisons use the same off-season of 2024 and are not independent replications. In 44 states with laboratory data, laboratory R_A correlated with hospital R_A at $r = 0.31$ and shared 0.230 (95% CI 0.043 to 0.418) of the hospital intercept.

4.9. Diminishing returns and surveillance design

Under $V(m) = a + b/m$, the share of the precision attainable at unlimited case volume is $m/(m + m_\times)$, and one doubling of m reduces variance by the fraction $m_\times/[2(m + m_\times)]$ (Proposition 1). The share is 33%, 50%, 67% and 80% at $m = 0.5$, 1, 2 and 4 times $m_\times$, a doubling from these points cuts variance by 33%, 25%, 17% and 10%, and 90% of attainable precision requires $m = 9m_\times$ (Fig. S5B). With the apparent balance ratios of 37.4 for R_A and 23.5 for G, 55% and 65% of state-weeks in the rule windows lay above the ratio of the corresponding statistic. These shares refer to the apparent scale. On the latent scale, 44% (95% CI 35 to 52) of state-weeks of R_A lay at or above the weekly balance scale of 58.9 admissions implied in the baseline setting. The share was 31% with a true mean generation interval of 2.9 days, 14% with 2.2 days and 57% with 4.3 days (Fig. 4C). For a single state, the number of observation weeks needed to reach a given relative standard error of mean R_A fell steeply with admissions below $m_\times$ and changed little above about $4m_\times$ (Fig. S5A). At the apparent balance ratio and a cost ratio of $c_0/c_1 = 10$, the optimal case volume was $m^* = 19.4$ (95% CI 13.8 to 26.6) for R_A and 15.3 for G; at $c_0/c_1 = 100$ it was 61.2 and 48.4.

The fitted variance departs from that of a Poisson renewal model, in which the error variance of the estimate is $\bar{R}/m$ and has no term that persists as m grows (Fig. S5D). On the apparent scale of R_A, the fitted slope was 1.25 times the Poisson value of 0.446, and the fitted variance was twice the Poisson variance at $m = 22$ and ten times at $m = 261$. On the latent scale of Proposition 1, with $v_R = 0.040$, the corresponding values were $m = 28$ and $m = 262$.

Seasons limit any estimate of variation between seasons (Fig. S6). Season-specific intercepts of R_A were 0.0137, 0.0062, 0.0127 and 0.0225, with a random-effects mean of 0.0125 (95% CI 0.0022 to 0.0227) and $I^2 = 0.47$ (Table S5). For G the random-effects mean was 0.0619 (95% CI 0.0066 to 0.1171) in rule windows and 0.0735 (95% CI -0.0558 to 0.2028) in calendar windows. The relative standard error of any unbiased estimate of the between-season variance is at least 1.00 with three seasons and 0.82 with four (Corollary 3). Intercept ratios for alternative denominator weights showed $q_1 < 1 < q_2$ for influenza, COVID-19 and RSV, for example $q_1 = 0.688$ and $q_2 = 1.539$ for influenza in calendar windows (Table S6). Epidemic growth with reporting noise and no environmental variation gave ranges of 0.63 to 0.89 for q_1 and 1.13 to 1.58 for q_2 in calendar windows, so this ordering does not discriminate between mechanisms. Of 40 one-sided tests, 21 were significant before correction, 18 after Benjamini–Hochberg correction and 10 after Holm correction, all of them for influenza and none for departures from the null ranges.

5. Discussion

The error of a case-count estimate of the reproduction number depends on the target, not only on the data. For the period-specific value, error falls without limit as reported cases accumulate. For the process mean it approaches the environmental variance, and the balance scale $m_\times = A/v_R$ marks the volume beyond which each doubling removes at most a quarter of the remaining variance. For the mean of the periods actually observed, the realised values are the target and the error is $A/(Tm)$. For the next period's value, no linear forecast from the current estimate falls below $(1 - \phi^2)v_R$ however many cases are counted (Corollary 1). In weekly influenza hospital admissions over four seasons, an affine model gave both apparent statistics positive intercepts, though an intercept-free power law fitted about as well; thinning supported the counting term, state intercepts roughly halved the apparent ratio within sampling variation, and the calibrated latent variance was about 0.04 at a 3.6-day generation interval.

As set out in the Introduction, the partition itself is classical, and the same form describes the variance of a cluster mean in two-stage sampling (Cochran, 1977). What the present model adds is its specification for case-count surveillance: negative binomial offspring enlarge the term that falls with case volume (Lloyd-Smith et al., 2005), binomial reporting adds to it, and reporting volatility adds a term independent of case volume that lowers the balance scale (Proposition 2). Because the target is averaged over periods, the balance scale also enters the division of a budget between cases per period and number of periods (Proposition 3), in the form of the optimal allocation in two-stage sampling and the optimal cluster size in cluster-randomised trials (Raudenbush, 1997). The contribution lies in this specification and in the calibration that links apparent weekly statistics to the latent variance.

The results complement the information-theoretic analyses of Parag and colleagues (Parag, 2021; Parag et al., 2022). Their metric for a noisy epidemic curve already contains time-varying reporting fractions, which enter through the geometric mean of the reporting probabilities, and a negative binomial extension is treated in their supplementary material (Parag et al., 2022). Their target is the sequence of period-specific values, and in a Poisson renewal model the error variance of each estimate is $\bar{R}/m$, so reporting acts by rescaling the effective case count. Here the target is the mean across periods, and variation shared by all chains of transmission in a period contributes variance independent of m. For the first target that variation is the signal; for the second it is sampling variation the target cannot shed. The two descriptions agree at small counts and separate near the balance scale. The Poisson analysis ranks streams by the information they carry about current transmission; the partition here indicates when further cases cease to improve estimates of typical transmission.

The thinning analysis tests one part of the measurement model directly in the data. Removing admissions at random leaves transmission, reporting volatility and sample composition unchanged, so only the term that falls with case volume should respond. The slope rose with $1/q$ at a rate compatible with the Poisson value (Fig. 3B), and the change in intercept was indistinguishable from zero at every level, although its intervals exclude only large changes. The slope at full retention exceeded the thinning rate, as expected if offspring overdispersion contributes a part of the slope that thinning does not amplify. The analysis therefore supports the A/m scaling of the counting term. It does not show that the intercept is environmental variation in transmission, because any source that does not scale with case volume, including reporting volatility and the construction of the statistic, would also leave the intercept unchanged.

The two-step inversion connects the apparent statistics with the latent model. The simulation first established that neither statistic estimates $v_R/\bar{R}^2$ directly; the intercept of G was 2.7 times this quantity and that of R_A about half of it, for reasons that follow from how each statistic aggregates growth over a week. After calibration the latent variances implied by the two statistics agreed within their sampling uncertainty (Section 4.6). That consistency is expected if both statistics respond to the same latent variation and the simulation captures how each transforms it, but as a test it has limited power: across the simulated settings it failed only when a synchronous shock to reported counts replaced the shock to transmission. The simulated panels also reproduced the intraclass correlation and the state-level lag-1 autocorrelation; the national-week share was set from the data and is not a test. Within the stated uncertainties the calibration is consistent with the empirical intercepts, under three conditions.

The first condition is the generation interval. The 3.6 days assumed by the statistics is the mean household serial

interval of Cowling et al. (2009), used as a proxy; the interval of 2.9 to 4.3 days is the confidence interval of that mean, not the range of plausible true intervals. Direct estimates of the mean generation interval, 2.8 to 3.2 days (Chan et al., 2025), and mean serial intervals of 2.2 and 2.8 days for influenza A subtypes (Vink et al., 2014) lie at or below the lower part of that interval. Across true mean intervals of 2.2 to 4.3 days the implied variance ranged from 0.013 to 0.058 for R_A. The inversion therefore identifies v_R conditionally on the interval, and its primary interval does not include uncertainty about it; a shorter true interval means a smaller latent variance and a larger latent balance scale.

The second condition is independence of weekly shocks. The pooled state-level autocorrelation lay near the centre of the simulated distribution under independence and below 97% of simulated panels with a first-order coefficient of 0.25 for R_A. A coefficient of 0.5 would have lowered the implied variance to about 0.027. This check concerns autocorrelation within states; the series of national-week means was not tested, and the simulated autocorrelation acted on the national and state shocks together.

The third condition concerns the cross-sectional slope. The calibrated theoretical b/a lies below both empirical point estimates, which fall at the 91st and 93rd percentiles of the simulated sampling distribution. Overdispersed admissions and a shorter generation interval can each raise the simulated ratio to the empirical values, and these explanations imply different latent variances. The state-level autocorrelation does not favour the shorter interval but cannot exclude it, and the data cannot exclude a mixture of transmission and reporting shocks (Proposition 2).

The fit across states also confounds case volume with spatial scale. States with more admissions are larger and aggregate more local epidemics; if local environmental variation is not fully synchronised within a state, Eq. (7) implies that the environmental variance of a state-level statistic falls with the number of local epidemics it contains, and part of the cross-sectional slope would then reflect aggregation and not counting. The fits with state intercepts are consistent with this explanation without establishing it: the within-state ratio was about half the cross-sectional ratio and lies inside the range the calibrated theory gives, and state intercepts tended to be lower in states with more admissions, but the paired ratio of the two values had an interval that included one. Spatial averaging therefore remains one of three candidate explanations of the high cross-sectional ratio, alongside overdispersed admissions and a shorter generation interval. If it operates, the cross-sectional ratio describes state-level aggregates of the size observed here, and the balance scale within a single reporting unit is lower.

Under the affine model, the empirical results are best read as a diminishing-returns scale for case volume within a period. A state already reporting at the apparent ratio obtains half of the precision that unlimited case volume could provide, and 90% requires nine times as many admissions. The limit applies to case volume per period: the error of an estimate pooled over T independent periods falls as v_R/T, and estimators that shrink towards a prior mean can go below the per-period value (Appendix B.5). A power law without an intercept described the observed range of scales about as well, so the balance scale is a statement about where returns diminish, not about error ceasing to fall with more data of every kind. The interpretation also does not depend on the source of the intercept: whatever mixture of transmission, testing and statistic construction it contains, a component shared within a period and not falling with case volume limits the value of further cases per period and raises the value of further periods.

The multiscale decomposition locates this component. About 41% of the intercept of R_A and 51% of that of G was common to all states in a calendar week, and resampling weeks moved the estimated intercept more than resampling states did. These shares refer to the intercept, not to the total variance of the statistic, of which the week component accounted for about a quarter to a third. A national-week component cannot be reduced by scaling up surveillance in a single state, because it enters every state's loss in that week; pooling states removes the state-specific part, and in these data pooled variance approached the national-week component once a few dozen states were pooled (Section 4.7). This share is consistent with national influences such as climate, school calendars and holidays (Cauchemez et al., 2008; Shaman et al., 2010), with national changes in testing or admission practice, and with a common trend within seasons; the present decomposition cannot separate these. Two-way inference is required for the same reason: intervals that resample states only were about two to five times too narrow for the intercept (Owen, 2007; Cameron et al., 2011).

The contrast between reporting regimes is descriptive. Under mandatory reporting a median of about 93% of hospitals reported, and adjustment for reporting completeness changed the intercept of R_A little, equivalently to zero within a margin of 17% (Fig. S4). Most state-weeks in the windows had nearly complete participation, so this bound concerns small departures from full participation, and at a reference volatility of 0.05 equivalence was not shown. During the voluntary period of 2024, when the median share fell to about 46%, the ratio of the interquartile range of log growth to that of the same months in 2023 was 1.58 (95% CI 0.95 to 2.23). The interval includes one, and the comparison rests on a single off-season; the placebo and matched comparisons preserve a ratio above one but are not independent replications. These results are compatible with more volatile growth statistics under low and variable facility participation, and they do not identify an effect of the reporting mandate. The adjustment measures facility participation only; testing intensity, coding and care seeking within reporting hospitals are not captured. The laboratory proxy shared a minority of the intercept but does not discriminate between transmission and ascertainment, because laboratory growth carries its own noise.

These results bear on the design of routine surveillance for the mean reproduction number. The threshold at which resources should shift from cases to periods is not the balance scale itself but the economic optimum $m^* = \sqrt{(c_0/c_1)m_\times}$, which combines the balance scale with the cost ratio. With independent periods, linear costs and a continuous number of periods, a system already collecting more than m^* cases per period reduces error by spending further resources on

periods and on jurisdictions whose variation is not synchronised (Eq. (7)); below m^* the same resources reduce error more when spent on cases. When R_t is autocorrelated, consecutive periods carry overlapping information, the optimum lay between a third of and 2.6 times m^* in the scenarios examined, and the threshold then requires a numerical optimisation (Appendix B).

The value of complete reporting depends on the reporting volatility a system actually has. In the hypothetical example of Section 4.2, with a relative volatility of $\tau = 0.15$, standardising reporting reduced error more than a tenfold increase in case volume per period. The completeness bound observed under mandatory reporting allows τ up to about 0.056 by the scaling against the reference volatility, or about 0.040 on the latent scale (Section 3.7), and at those volatilities case volume remained the more effective lever. The value of complete reporting relative to case volume grows with the volatility that incomplete participation introduces, which the voluntary period suggests may be larger but did not establish. Seasons are the limiting unit for statements about variation between seasons: with four seasons no unbiased estimator of a between-season variance can have a relative standard error below 0.82 (Corollary 3), and the intercept ratios for alternative denominator weights did not discriminate between mechanisms, because epidemic growth with reporting noise alone already produced $q_1 < 1 < q_2$.

The scope of these conclusions is set by the data, the model form and the calibration. The empirical results rest on one surveillance system, four influenza seasons and rule-defined early-growth windows; the window rule was revised once and the renewal-type ratio became the primary statistic during the analysis, and Table S9 lists each specification with its role and date. The positive intercept and every quantity derived from it are conditional on the affine form, which the data did not prefer over a power law. Thinning tests only the Poisson scaling of the counting term and leaves the source of the intercept unidentified. The fits with state intercepts do not separate state and week effects. The intercept pools all seasons around one mean. The inversion is conditional on the generation interval and on the simulated delay, offspring and independence structure. The statistic R_A is a transformation of weekly growth, not a posterior estimate of R_t (Cori et al., 2013). The propositions assume environmental variation independent between periods unless stated otherwise, a case volume independent of R_t and a fixed, known reporting probability; the renewal formulation of Section 3.3 shows that uncertainty in infection pressure adds a further term that does not fall with case volume.

Three extensions would address the main conditions directly: an inversion that integrates over uncertainty in the generation interval, finer spatial units within states to test the spatial explanation of the cross-sectional ratio, and simulations with autocorrelation in the national shock alone and with the window rule applied to simulated data. Renewal estimates with an explicit admission delay, hierarchical models with week effects and independent measures of testing, and a second surveillance system with mandatory reporting and more seasons would extend the empirical base.

6. Conclusion

The value of further reported cases for estimating transmission depends on the inferential target. For the current value, error falls without limit as cases accumulate. For the process mean across periods, between-period variation of R_t is sampling variation of the target itself, so precision grows with case volume at a diminishing rate; beyond the balance scale each further doubling of cases removes at most a quarter of the remaining variance. For the next period's value, no linear forecast from the current estimate has error below $(1 - \phi^2)v_R$.

In US influenza hospital admissions, an affine model placed the apparent balance ratio near 37 admissions per state-week, although a power law without a floor fitted about as well. Thinning steepened the slope as Poisson counting predicts and left the intercept unchanged; state intercepts lowered the ratio to about 17, within sampling variation. Calibration implied an environmental variance near 0.04 at a 3.6-day generation interval and 0.013 to 0.026 at 2.2 to 2.9 days, and a substantial share of the intercept was common to all states in a week.

With independent periods and linear costs, a system above $\sqrt{(c_0/c_1)m_\times}$ cases per period gains more from further periods and weakly synchronised regions than from further cases. The reporting evidence is narrower: under mandatory reporting facility completeness changed the intercept little, and the 2024 voluntary comparison was inconclusive. Empirical intercepts of apparent statistics should be reported with the calibration that links them to transmission and with the functional form they depend on.

Data and code availability

The NHSN Weekly Hospital Respiratory Data are publicly available from the U.S. Centers for Disease Control and Prevention (Centers for Disease Control and Prevention, 2026); the analyses used the version retrieved on 2026-10-20. The code for the simulations, estimation and figures is available at https://github.com/wangchi2068/case-count-bottleneck.

A. Proofs

A.1. Offspring and reporting model

The negative binomial distribution $\mathrm{NB}(r, p)$ with r successes and success probability p has mean $r(1-p)/p$ and variance $r(1-p)/p^2$. Setting $r = k$ and $p = k/(k + R_t)$ gives Eq. (1). A sum of n independent $\mathrm{NB}(k, p)$ variables is $\mathrm{NB}(nk, p)$, and binomial thinning of $\mathrm{NB}(r, k/(k + \mu))$ with retention probability ρ gives $\mathrm{NB}(r, k/(k + \rho\mu))$, which together give Eq. (2).

A.2. The period-specific target

Write c for an observed value of C_t. From Eq. (2), the log-likelihood of R_t given n is given by:

$$\ell(R_t) = \text{const} + nk \log \frac{k}{k + \rho R_t} + c \log \frac{\rho R_t}{k + \rho R_t}, \tag{A.1}$$

and the score can be written as:

$$\dot{\ell}(R_t) = \frac{c}{R_t} - \frac{(c + nk)\rho}{k + \rho R_t} = \mathcal{J}(R_t)\,(\hat{R}_t - R_t), \qquad \mathcal{J}(R_t) = \frac{n}{R_t/\rho + {R_t}^2/k}, \tag{A.2}$$

where $\mathcal{J}(R_t)$ is the Fisher information. The conditional mean of $\hat{R}_t = c/(\rho n)$ is R_t, and its conditional variance is $R_t/(\rho n) + {R_t}^2/(kn) = 1/\mathcal{J}(R_t)$. Because the score in Eq. (A.2) is proportional to $\hat{R}_t - R_t$, the Cramér–Rao bound is attained for every finite n. The distribution in Eq. (2) is a one-parameter exponential family in C_t with natural parameter $\log[\rho R_t/(k + \rho R_t)]$, which ranges over an open interval, so C_t is complete and sufficient. By the Lehmann–Scheffé theorem, $\hat{R}_t$ is the uniformly minimum-variance unbiased estimator of R_t (Lehmann and Casella, 1998; Lehmann and Romano, 2005).

A.3. Proof of Proposition 1

Under H4, $\mathbb{E}(R_t \mid n) = \bar{R}$ and $\mathrm{Var}(R_t \mid n) = v_R$. By the law of total variance and Appendix A.2, the error variance is given by:

$$\mathrm{Var}(\hat{R}_t \mid n) = \mathbb{E}\left[\frac{R_t}{\rho n} + \frac{{R_t}^2}{kn}\right] + \mathrm{Var}(R_t) = v_R + \frac{\bar{R}}{\rho n} + \frac{\bar{R}^2 + v_R}{kn}, \tag{A.3}$$

and substituting $m = \rho n$ gives Eq. (4). Parts (i) and (ii) follow by taking limits. For part (iii), $[V(m) - V(2m)]/V(m) = [A/(2m)]/(v_R + A/m) = 1/[2(1 + m/m_\times)]$, which equals one quarter at $m = m_\times$. For part (iv), let $\tilde{\theta}$ be any estimator from data on period t with $\mathbb{E}(\tilde{\theta} \mid R_t) = R_t$. Then $\mathbb{E}(\tilde{\theta} - \bar{R})^2 = \mathbb{E}[\mathrm{Var}(\tilde{\theta} \mid R_t)] + v_R \geq \mathbb{E}[1/\mathcal{J}(R_t)] + v_R = V$, with equality for $\hat{R}_t$. The first equality uses the conditional unbiasedness of $\tilde{\theta}$, which makes $\tilde{\theta} - R_t$ uncorrelated with $R_t - \bar{R}$; the same argument applied to $\hat{R}_t$ gives the difference of exactly v_R between the mean squared errors about $\bar{R}$ and about R_t. For prediction of the next scaled count $y = C_{t+1}/m$ when R_{t+1} is independent of the history, $\mathbb{E}(\hat{y} - y)^2 = (\hat{y} - \mathbb{E}y)^2 + \mathrm{Var}(y) \geq V$, with equality at $\hat{y} = \bar{R}$.

The marginal effects of coverage and completeness are $-\partial V/\partial \log n = A/m$ and $-\partial V/\partial \log \rho = \bar{R}/m$, so their ratio is $A/\bar{R}$. If R_t follows a stationary first-order autoregressive process with variance v_R and autocorrelation ϕ, the best linear predictor of R_{t+1} given R_t itself has error variance $(1 - \phi^2)v_R$ (Hamilton, 1994); a predictor that sees only $\hat{R}_t$ cannot do better (Appendix A.6). For J jurisdictions with common variance v_R and mean pairwise correlation $\bar{c}$, the variance of their average is $J^{-2}[Jv_R + J(J-1)\bar{c}v_R]$, which gives Eq. (7).

A.4. Proof of Proposition 2

Write $\sigma^2 = \mathrm{Var}(p_t) = \tau^2\rho^2$. Given (n, R_t, p_t), C_t has mean $p_t n R_t$ and variance $p_t n R_t + {p_t}^2 n {R_t}^2/k$. Averaging over p_t, the conditional variance given (n, R_t) is given by:

$$\mathrm{Var}(C_t \mid n, R_t) = \rho n R_t + \frac{(\rho^2 + \sigma^2) n {R_t}^2}{k} + \sigma^2 n^2 {R_t}^2. \tag{A.4}$$

The conditional mean is $\rho n R_t$, so $\hat{R}_t$ remains unbiased for $\bar{R}$. Averaging over R_t and dividing by $(\rho n)^2$ gives:

$$V_H(m) = v_R + \tau^2(\bar{R}^2 + v_R) + \frac{\bar{R}}{m} + \frac{\rho(1 + \tau^2)(\bar{R}^2 + v_R)}{km}, \tag{A.5}$$

which is Eq. (8). Writing $U = (1+\tau^2)(\bar{R}^2+v_R)$, the limiting error variance is $v_R+v_H = U-\bar{R}^2$ and $A_H = \bar{R}+\rho U/k$, so the mean and variance of $\hat{R}_t$ depend on v_R and τ only through U. With $Q = \bar{R}^2+v_R$, $v_H = \tau^2 Q$ and $A_H - A = \rho\tau^2 Q/k$, the difference of the two balance scales is:

$$m_\times - m_{\times,H} = \frac{A(v_R + v_H) - A_H v_R}{v_R(v_R + v_H)} = \frac{\tau^2 Q\,[A - \rho v_R/k]}{v_R(v_R + v_H)} = \frac{\tau^2 Q\,\bar{R}\,(1 + \rho\bar{R}/k)}{v_R(v_R + v_H)}, \tag{A.6}$$

because $A - \rho v_R/k = \bar{R} + \rho\bar{R}^2/k$. The right-hand side is zero at $\tau = 0$ and positive for $\tau > 0$, which gives $m_{\times,H} \leq m_\times$ with equality if and only if $\tau = 0$, and Eq. (9).

A.5. *Proof of Proposition 3*

Conditional on $(R_1, \dots, R_T)$, counts in different periods are independent. With $\bar{R}_T = T^{-1}\sum_t R_t$, the variance of the mean estimate is given by:

$$\mathrm{Var}\Big(\frac{1}{T}\sum_{t=1}^{T}\hat{R}_t\Big) = \mathrm{Var}(\bar{R}_T) + \frac{A}{Tm}, \tag{A.7}$$

where $\mathrm{Var}(\bar{R}_T) = v_R/T$ for independent periods, giving Eq. (10). Substituting $T = \mathcal{B}/(c_0 + c_1 m)$ into Eq. (10) gives the objective:

$$f(m) = \frac{(c_0 + c_1 m)(v_R + A/m)}{\mathcal{B}} = \frac{c_0 v_R + c_1 A + c_1 v_R m + c_0 A/m}{\mathcal{B}}, \tag{A.8}$$

which is strictly convex for $c_0 > 0$, with $f'(m) = (c_1 v_R - c_0 A/m^2)/\mathcal{B}$. The unique stationary point is $m^{*2} = c_0 A/(c_1 v_R) = (c_0/c_1) m_\times$, which is Eq. (11). Because f is strictly convex, $f'(m) > 0$ if and only if $m > m^*$, so at a fixed budget the error increases with m exactly when $m > m^*$.

A.6. *Proof of Corollary 1*

Write $e_t = \hat{R}_t - R_t$. Given (R_s), the e_t are independent with mean zero and $\mathbb{E}({e_t}^2 \mid R_t) = R_t/m + \rho {R_t}^2/(km)$, so $\mathbb{E}{e_t}^2 = A/m$, and e_t is uncorrelated with every function of (R_s). For part (a), $T^{-1}\sum_t \hat{R}_t - \bar{R} = (\bar{R}_T - \bar{R}) + T^{-1}\sum_t e_t$, and the two terms are uncorrelated, so the error is $\mathrm{Var}(\bar{R}_T) + A/(Tm)$. For a stationary first-order autoregressive process with $\mathrm{Cov}(R_s, R_t) = v_R \phi^{|s-t|}$, $T^2\mathrm{Var}(\bar{R}_T) = v_R[T + 2\sum_{h=1}^{T-1}(T-h)\phi^h]$, and summing the geometric series gives Eq. (5) (Hamilton, 1994). For part (b), $T^{-1}\sum_t \hat{R}_t - \bar{R}_T = T^{-1}\sum_t e_t$, whose mean square is $A/(Tm)$ for any law of (R_t). For part (c), $\hat{R}_t - R_{t+1} = e_t + (R_t - R_{t+1})$, and $\mathbb{E}(R_t - R_{t+1})^2 = 2(1-\phi)v_R$, which gives the naive error. For a forecast $\bar{R} + \lambda(\hat{R}_t - \bar{R})$, $\mathrm{Cov}(\hat{R}_t, R_{t+1}) = \phi v_R$ and $\mathrm{Var}(\hat{R}_t) = v_R + A/m$, so the error is:

$$v_R - 2\lambda\phi v_R + \lambda^2\left(v_R + \frac{A}{m}\right), \tag{A.9}$$

which is minimised at $\lambda = \phi v_R/(v_R + A/m)$ with value $v_R - \phi^2 {v_R}^2/(v_R + A/m)$. Since $v_R/(v_R + A/m) = m/(m + m_\times)$, this is Eq. (6). The value decreases in m towards $(1-\phi^2)v_R$, so no linear forecast from $\hat{R}_t$ has a smaller error.

A.7. *Proof of Corollary 2*

Write $U_0 = \bar{R}^2 + v_R$. From Eq. (8), the gains from proportional increases in coverage and completeness are given by:

$$-\frac{\partial V_H}{\partial \log n} = \frac{\bar{R}}{\rho n} + \frac{(1+\tau^2)U_0}{kn}, \qquad -\frac{\partial V_H}{\partial \log \rho} = \frac{\bar{R}}{\rho n} - \rho\frac{\mathrm{d}\tau^2}{\mathrm{d}\rho}\,U_0\left(1 + \frac{1}{kn}\right). \tag{A.10}$$

If $\mathrm{d}\tau^2/\mathrm{d}\rho = 0$, the first gain exceeds the second by $(1+\tau^2)U_0/(kn) > 0$. Under binomial sampling of ν units, $\mathrm{Var}(p_t) = \rho(1-\rho)/\nu$, $\tau^2 = (1-\rho)/(\rho\nu)$ and $\rho\,\mathrm{d}\tau^2/\mathrm{d}\rho = -1/(\rho\nu)$. The second gain minus the first is

$$\frac{U_0}{kn}\left[\frac{kn + \rho}{\rho\nu} - 1\right], \tag{A.11}$$

which is positive if and only if $kn > \rho\nu - \rho$, that is $m = \rho n > \rho^2(\nu - 1)/k$. Under a Beta prior $\mathrm{Var}(p_t) = \rho(1-\rho)/(\nu+1)$, the denominator inside the brackets becomes $\rho(\nu+1)$ and the condition is $m > \rho^2\nu/k$.

A.8. *Renewal formulation and denominator noise*

Let $\bar{\Lambda} = \mathbb{E}(\Lambda_{t+1} \mid \mathcal{F}_t)$ and $m = \rho\bar{\Lambda}$. Assuming that given (Λ_{t+1}, R_{t+1}), incident cases follow a negative binomial distribution, and given $\mathcal{F}_t$, Λ_{t+1} and R_{t+1} are independent with $\mathbb{E}(R_{t+1} \mid \mathcal{F}_t) = \bar{R}$ and $\mathrm{Var}(R_{t+1} \mid \mathcal{F}_t) = v_R$, conditioning on (Λ_{t+1}, R_{t+1}) and applying the law of total variance gives:

$$\mathrm{Var}(C_{t+1} \mid \mathcal{F}_t) = mA + \rho^2\bar{\Lambda}^2\left[v_R + \mathrm{CV}^2(\Lambda_{t+1} \mid \mathcal{F}_t)\,(\bar{R}^2 + v_R)\right], \tag{A.12}$$

and dividing by m^2 gives Eq. (15). For Eq. (17), assuming a flat latent trajectory ($Z_{t-j} \equiv Z$) with independent lagged reporting noise, write $D_t = m(1+\delta)$ with $\mathbb{E}\delta = 0$ and $\mathbb{E}\delta^2 = c_w/m$, where $c_w = (1-\rho)\sum_s {w_s}^2$, and let $Y = C_{t+1}/m$ be independent of δ with mean $\bar{R}$ and variance $v_R + A/m$. Expanding $G_t = Y(1+\delta)^{-1}$ to second order gives:

$$\mathbb{E}(G_t) = \bar{R}\left(1 + \frac{c_w}{m}\right) + O(m^{-2}), \qquad \mathrm{Var}(G_t) = v_R + \frac{A + (\bar{R}^2 + 3v_R)c_w}{m} + O(m^{-2}), \tag{A.13}$$

and dividing by $[\mathbb{E}(G_t)]^2 = \bar{R}^2(1 + 2c_w/m) + O(m^{-2})$ gives a relative variance of $v_R/\bar{R}^2 + [A + (\bar{R}^2 + v_R)c_w]/(\bar{R}^2 m)$ to first order in $1/m$, which is Eq. (17).

B. Extensions

B.1. Precision of the environmental variance

Proposition 4 (Precision of environmental variance)**.** *Suppose that $R_1, \dots, R_T$ are independent gamma variables with known mean $\bar{R}$ and variance v_R, and let $\xi = \bar{R}^2/v_R$. Any unbiased estimator of v_R based on the counts from the T periods has relative standard error at least*

$$\left[T\xi\left(\xi\psi'(\xi) - 1\right)\right]^{-1/2}, \tag{B.1}$$

where ψ' is the trigamma function. As $\xi \to \infty$, the bound tends to the Gaussian value $\sqrt{2/T}$.

Proof. The counts depend on v_R only through $R_{1:T}$, so by the data processing inequality the Fisher information for v_R in the counts cannot exceed that in $R_{1:T}$. With the mean fixed, the gamma shape is ξ and the scale is $v_R/\bar{R}$. The information for ξ in one observation is $\psi'(\xi) - 1/\xi$, and $\mathrm{d}\xi/\mathrm{d}v_R = -\xi/v_R$, so the information for v_R in T periods is $T\xi(\xi\psi'(\xi) - 1)/{v_R}^2$. The Cramér–Rao inequality gives Eq. (B.1). The limit follows from $\psi'(\xi) = 1/\xi + 1/(2\xi^2) + O(\xi^{-3})$. □

At the baseline of Table S7 ($\bar{R} = 1.15$, $v_R = 0.04$), $\xi = 33.06$ and $\xi(\xi\psi'(\xi) - 1) = 0.505$, so the relative standard error is at least 0.812 with 3 periods, 0.445 with 10 periods and 0.199 with 50 periods. The constant 0.505 is specific to this baseline, and Eq. (B.1) is the general form.

B.2. Budget threshold with autocorrelated environments

With autocorrelation ϕ, the environmental term of Proposition 3 is $v_R e(T)$ with $e(T) = S(\phi, T)/T^2$. Define $g_T = -T^2 e'(T)$, treating T as continuous, so that $g_T = 1$ for independent periods, and let ${m_0}^* = \sqrt{(c_0/c_1)m_\times}$ be the optimum of Eq. (11).

Proposition 5 (Budget threshold)**.** *Under the budget constraint $\mathcal{B} = T(c_0 + c_1 m)$, the optimal case volume per period satisfies $m^* = {m_0}^*(g_{T^*})^{-1/2}$, where T^* is the optimal number of periods. If g_T increases in T from $g_1 < 1$ towards $(1+\phi)/(1-\phi) > 1$, there is a unique $T^\dagger$ with $g_{T^\dagger} = 1$, and m^* exceeds ${m_0}^*$ when $T^* < T^\dagger$ and falls below it when $T^* > T^\dagger$. The two regimes meet at the budget*

$$\mathcal{B}^\dagger = T^\dagger\left(c_0 + c_1 {m_0}^*\right), \tag{B.2}$$

which is expressed in units of the cost per unit case volume when $c_1 = 1$.

Proof. Writing the sampling term as $Ac_1/(\mathcal{B} - c_0 T)$ and differentiating $v_R e(T) + Ac_1/(\mathcal{B} - c_0 T)$ with respect to T gives the first-order condition $v_R g_T = c_0 A/(c_1 m^2)$, so ${m^*}^2 = {{m_0}^*}^2/g_{T^*}$. The ordering of m^* and ${m_0}^*$ follows from the monotonicity of g_T, and at $T^* = T^\dagger$ the budget constraint gives Eq. (B.2). □

For $\phi = 0.8$, g_T increases in T and $T^\dagger = 4.484$. With $c_0/c_1 = 100$ and the baseline balance scale, ${m_0}^* = 72.86$ and $\mathcal{B}^\dagger = 775.0$; over the uncorrected scenario range of 21.3 to 49.2 cases per generation in Section 3.3, $\mathcal{B}^\dagger$ ranges from 655.3 to 762.9 (762.85 rounded, Table S7). The calculation treats T as continuous and the environmental process as stationary, so it applies within a season and not across seasons with different mean transmission.

B.3. Between-season variance

Corollary 3 (Season-level bound)**.** *Suppose that season effects $\eta_1, \dots, \eta_S$ are independent Gaussian variables with unknown mean and variance v_S, and that the counts depend on v_S only through the season effects. Any unbiased estimator of v_S based on the counts has relative standard error at least $\sqrt{2/(S-1)}$, which is 1.00 for $S = 3$ and 0.816 for $S = 4$.*

Proof. Let δ be an unbiased estimator based on the counts. Because the conditional distribution of the counts given $\eta_{1:S}$ does not depend on the unknown parameters (μ, v_S) and unbiasedness holds for all (μ, v_S), $\mathbb{E}(\delta \mid \eta_{1:S})$ is a function of the season effects that is unbiased for v_S, and its variance does not exceed that of δ. The sample mean and sample variance of $\eta_{1:S}$ are complete and sufficient, so the sample variance is the minimum-variance unbiased estimator of v_S, with variance $2{v_S}^2/(S-1)$ (Lehmann and Casella, 1998; Lehmann and Romano, 2005; Searle et al., 1992). □

The bound uses $S-1$ degrees of freedom because the mean across seasons is estimated; the corresponding bound with a known mean, $\sqrt{2/S}$, is smaller (Table S7). Season-specific intercepts for three and four seasons are reported separately in Table S5.

B.4. Testing whether transmission exceeds one

For $R > R'$, the likelihood ratio of Eq. (2) increases in c because $R/(k+\rho R)$ increases in R, so the family has a monotone likelihood ratio in C_t. The test that rejects $R_t \le 1$ when $C_t \ge c_\alpha$ therefore has a power function that does not decrease in R_t, and by the Karlin–Rubin theorem it is uniformly most powerful at its attained level $\alpha^* = P_{R_t=1}(C_t \ge c_\alpha) \le \alpha$ (Karlin and Rubin, 1956; Lehmann and Romano, 2005). Because C_t is discrete, α^* is generally below the nominal level; optimality holds at α^*, not at α, unless the test is randomised or a mid-p value is used.

B.5. Shrinkage as an ideal bound

Part (iv) of Proposition 1 excludes estimators that are biased given R_t. Consider $\hat{\theta}_\lambda = (1-\lambda)\hat{R}_t + \lambda\mu_0$ for a fixed benchmark μ_0 with $\Delta = \mu_0 - \bar{R}$. Its mean squared error for $\bar{R}$ is given by:

$$\text{MSE}(\lambda) = (1-\lambda)^2 V + \lambda^2 \Delta^2, \tag{B.3}$$

where V is given by Eq. (4). The minimum is at $\lambda^* = V/(V+\Delta^2)$, where $\text{MSE}(\lambda^*) = V\Delta^2/(V+\Delta^2)$, which is below v_R if and only if $\Delta^2 < v_R(1 + m/m_\times)$. The weight λ^* depends on the unknown $\bar{R}$ and V, so this is an ideal bound and not a feasible estimator; an empirical Bayes or hierarchical estimator estimates the weight from the data and does not attain it (Stoner et al., 2019).

C. Supplementary figures

Figures S1–S6 give the detailed panels behind the main figures. They show the budget example and the scenario range of the optimal case volume (Fig. S1) and the calibration of the inversion (Fig. S2). They also show the national series and sensitivity of the empirical fit (Fig. S3), the reporting regimes (Fig. S4), the design nomogram (Fig. S5) and the between-season analysis (Fig. S6).

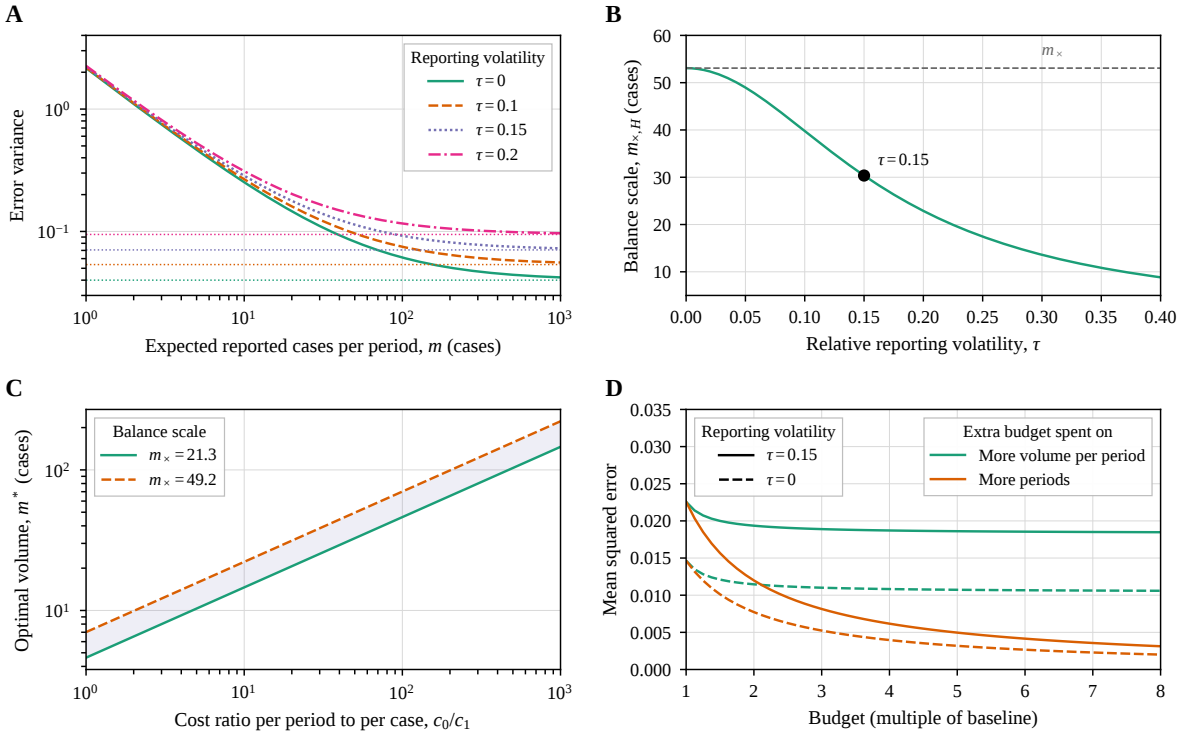


Fig. S1. Reporting heterogeneity and the allocation of a fixed surveillance budget. (A) Error variance against expected reported cases per period for four levels of relative reporting volatility; dotted lines mark the corresponding limits as case volume grows, set by environmental plus reporting variance. (B) Generalised balance scale against relative reporting volatility; the dashed line is the balance scale without reporting volatility, and the dot marks the volatility used in panel D. (C) Optimal number of reported cases per period against the price ratio c_0/c_1, for the two ends of the uncorrected scenario range of the balance scale, 21.3 and 49.2 cases per generation (labelled $m_\times$ in the panel); the band spans the range between them. (D) Mean squared error of the estimator of the mean reproduction number against total budget, when additional budget buys either more cases per period or more periods, with autocorrelated environmental and reporting variation, under baseline (solid) and standardised (dashed) reporting.

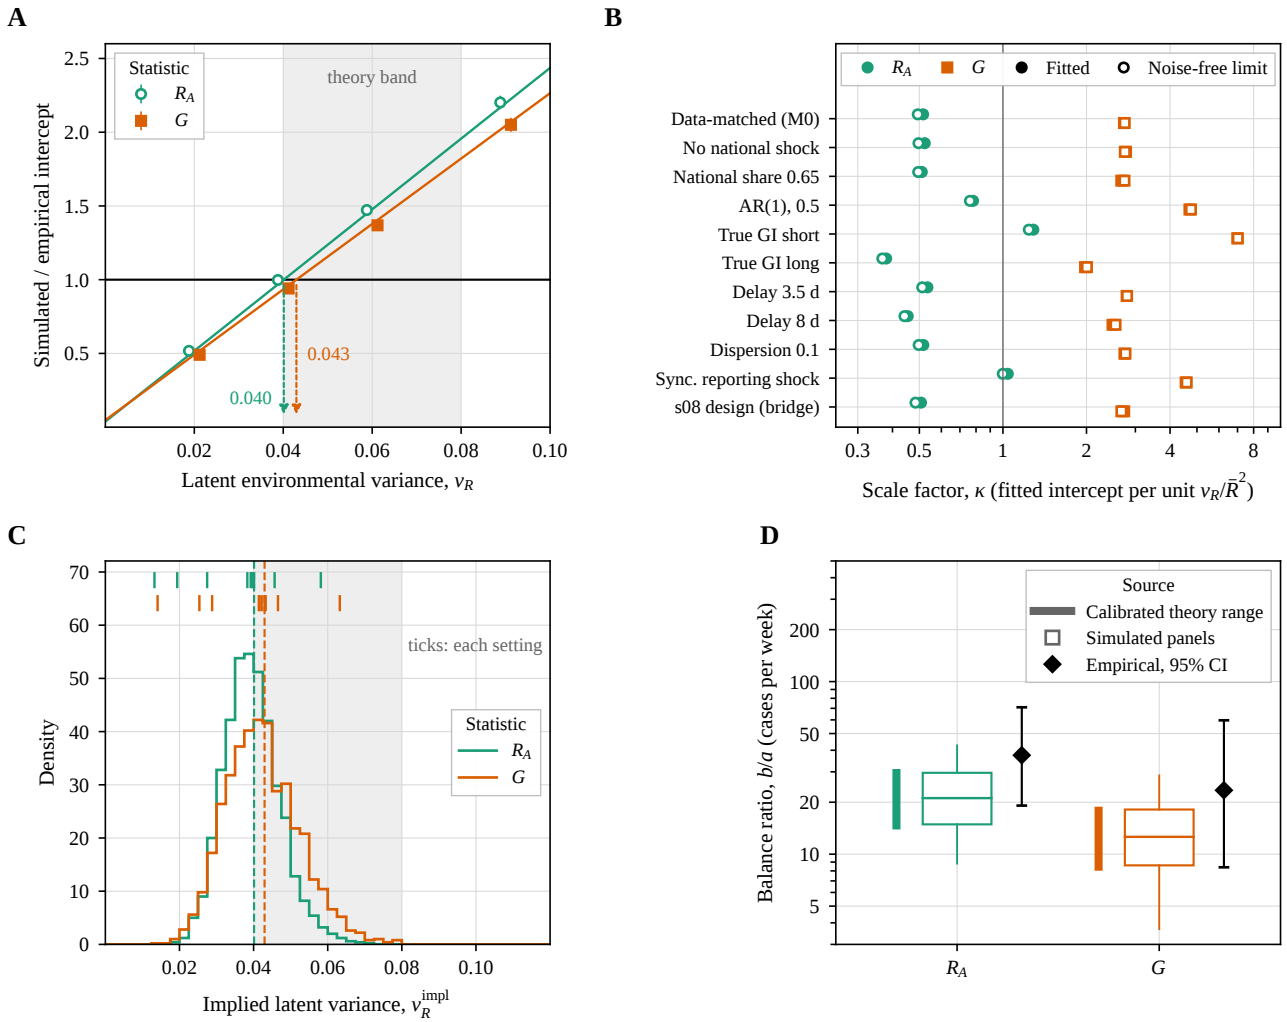


Fig. S2. Simulation-based inversion from apparent statistics to latent dynamics. (A) Calibration lines from 200 simulated panels per cell that reproduce the observed seasons, window lengths, state sizes, national week shocks and reporting delay. Fitted intercepts are scaled by the empirical intercept. Dashed drop lines mark where each line meets the data, giving the implied latent variance v_R. (B) Scale factor κ linking the fitted intercept to $v_R/\bar{R}^2$ for each simulation setting. Filled markers are fitted values and open markers the noise-free limit. (C) Distribution of the implied v_R under the data-matched setting. Uncertainty combines the two-way bootstrap of the data with Monte Carlo error of the calibration. Top ticks show the point values of the other settings. (D) Simulated per-panel b/a and the calibrated theoretical range, against empirical estimates with 95% intervals. Grey bands mark $v_R = 0.04$–0.08.

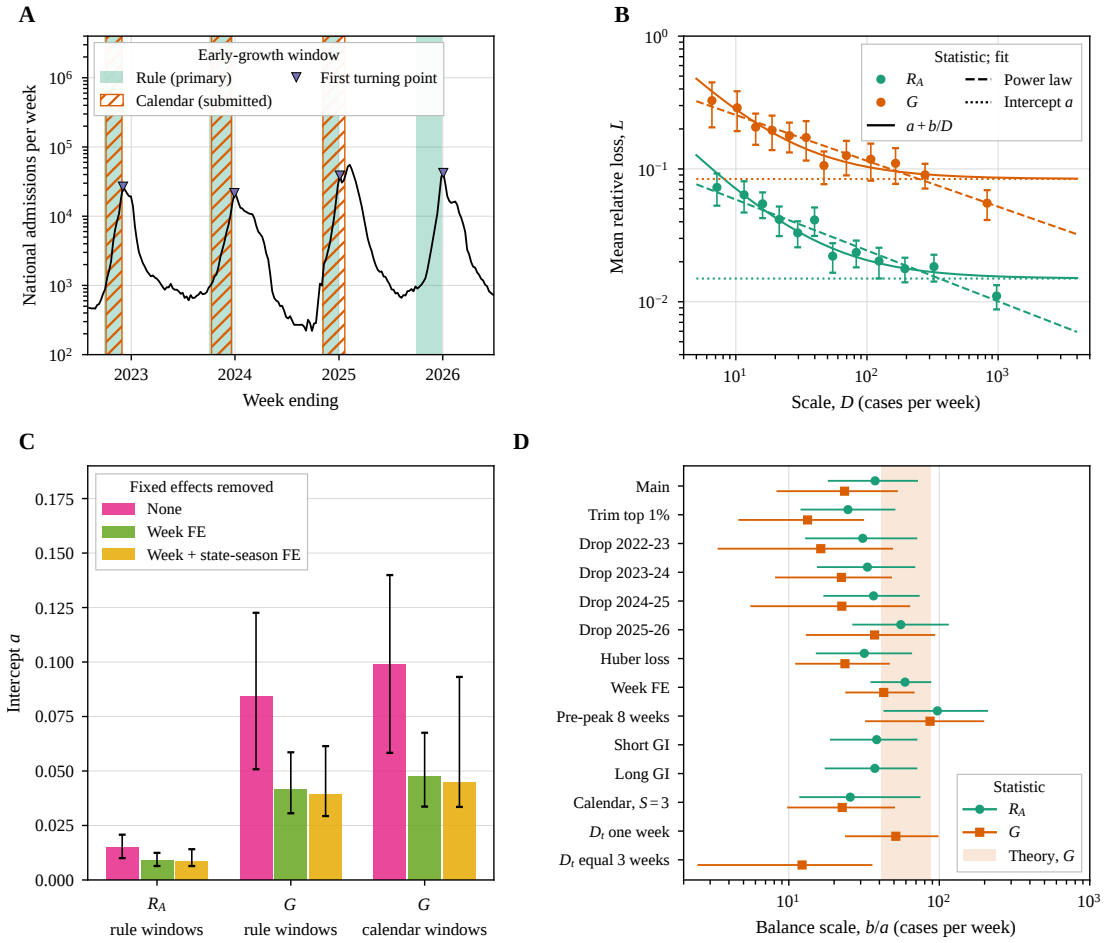


Fig. S3. Relative loss of two apparent weekly statistics in influenza hospital admissions. (A) National weekly admissions on a logarithmic scale; shading marks the rule-based early-growth windows of the primary analysis, hatching marks the calendar windows of the earlier analysis, and triangles mark the first national turning point of each season. (B) Binned mean relative loss against case-count scale for the renewal-type ratio and the weighted growth ratio, with the fitted affine curve, a power law without intercept used for comparison and the fitted intercept; bin error bars ignore clustering. (C) Intercept before adjustment, after removing national-week effects and after also removing state-season effects, with two-way bootstrap intervals. (D) Ratio of slope to intercept under sensitivity variants; the band is the uncalibrated theoretical range for the growth ratio with the denominator correction, and Fig. S2D gives the calibrated range.

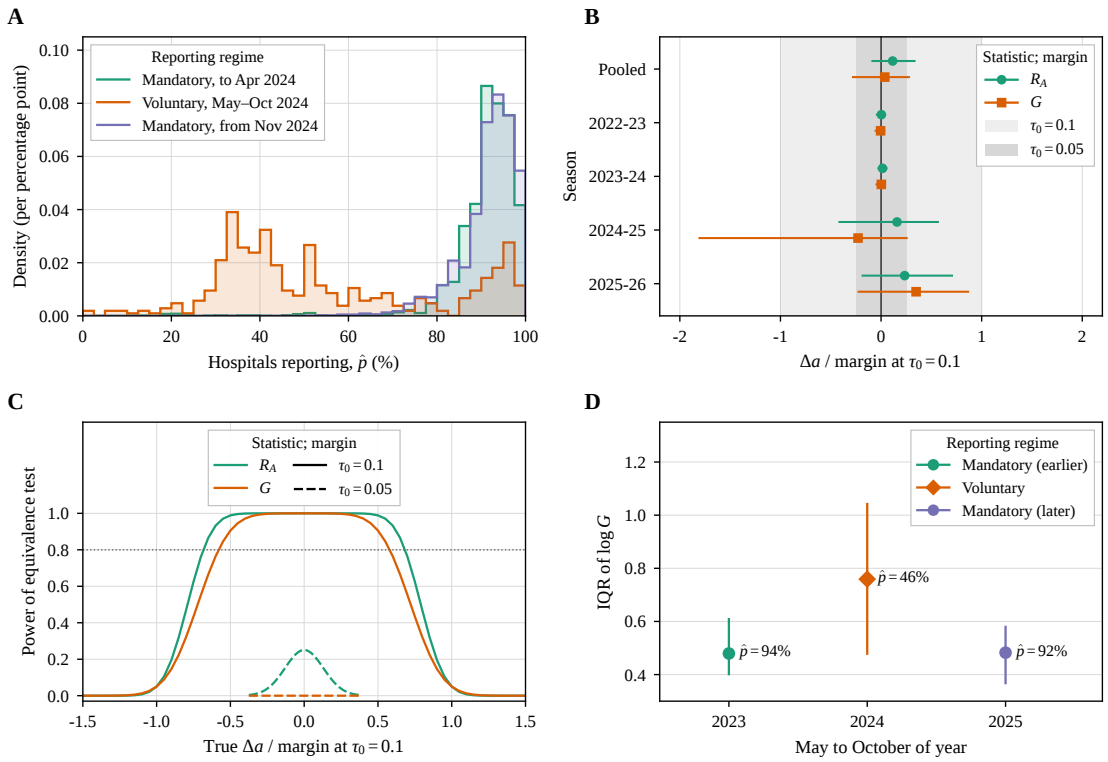


Fig. S4. Reporting completeness and the empirical intercept. (A) Distribution of the share of hospitals reporting influenza admissions under the three reporting regimes. (B) Change in the intercept when the statistic is adjusted for reporting coverage, pooled and by season, with 90% paired bootstrap intervals. Each statistic is expressed relative to its own equivalence margin, and shaded bands show the margins implied by two thresholds of reporting volatility. (C) Power of the two one-sided equivalence test against the true change, for both statistics and both margins, with a conventional power reference. (D) Interquartile range of the logarithm of the growth ratio over May to October in three calendar years; the year with voluntary reporting is highlighted, and labels give median reporting completeness.

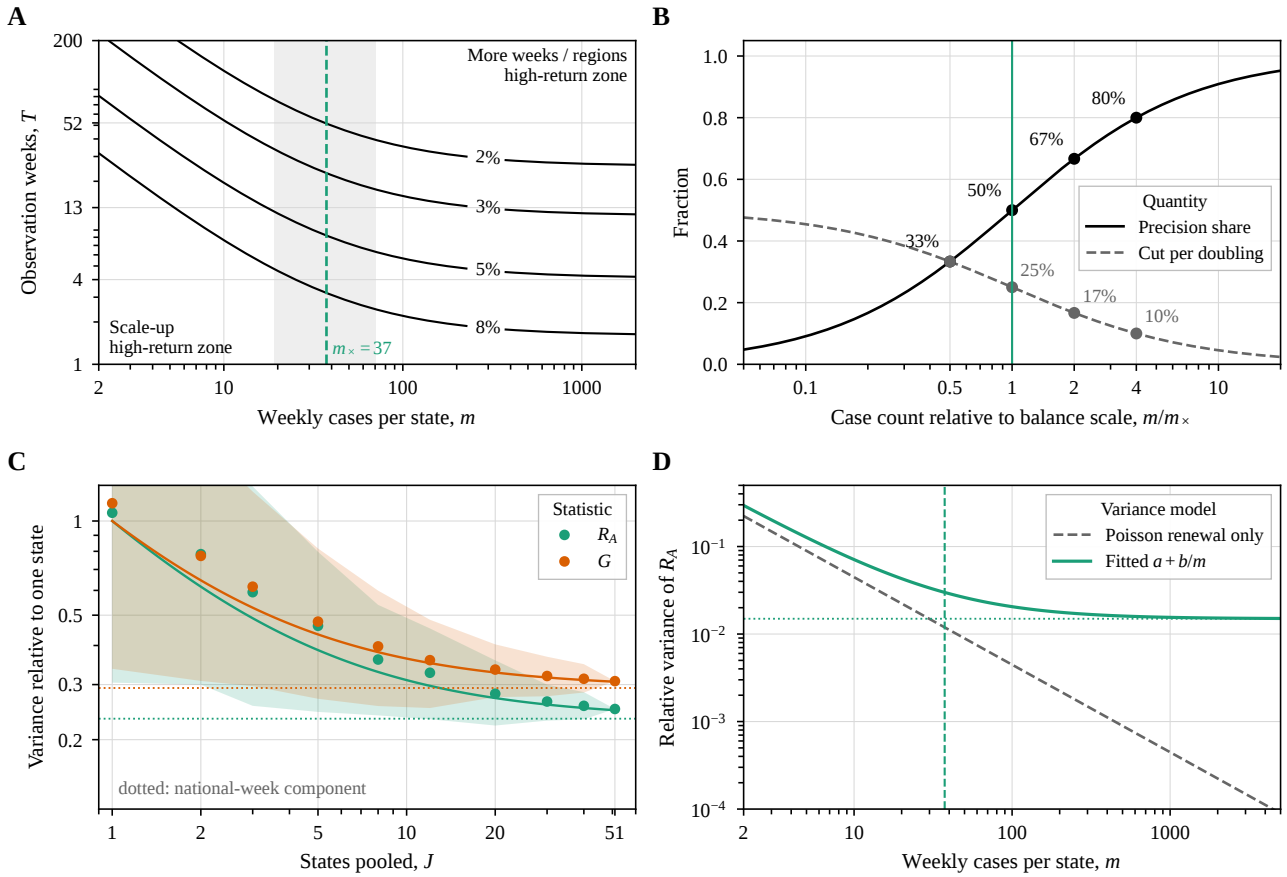


Fig. S5. Surveillance design nomogram. (A) Observation weeks needed for the relative standard error of a single state's mean renewal ratio R_A to reach 2–8%, by weekly admissions m, on the scale of the statistic and ignoring autocorrelation within states. The dashed line marks the apparent balance ratio b/a of R_A, labelled $m_\times$ in the panel, and the grey band its 95% interval. To its left, added cases give most of the gain. To its right, contours flatten and added weeks or regions give most of the gain. (B) Share of attainable precision, $m/(m + m_\times)$, and the variance reduction from one doubling of m. (C) Variance of the mean weekly deviation pooled over J states, relative to one state, for R_A and the growth ratio G. Points and bands come from random state subsets; dotted lines show the national-week component. (D) Relative variance of R_A under the Poisson renewal variance alone and under the fitted $a + b/m$.

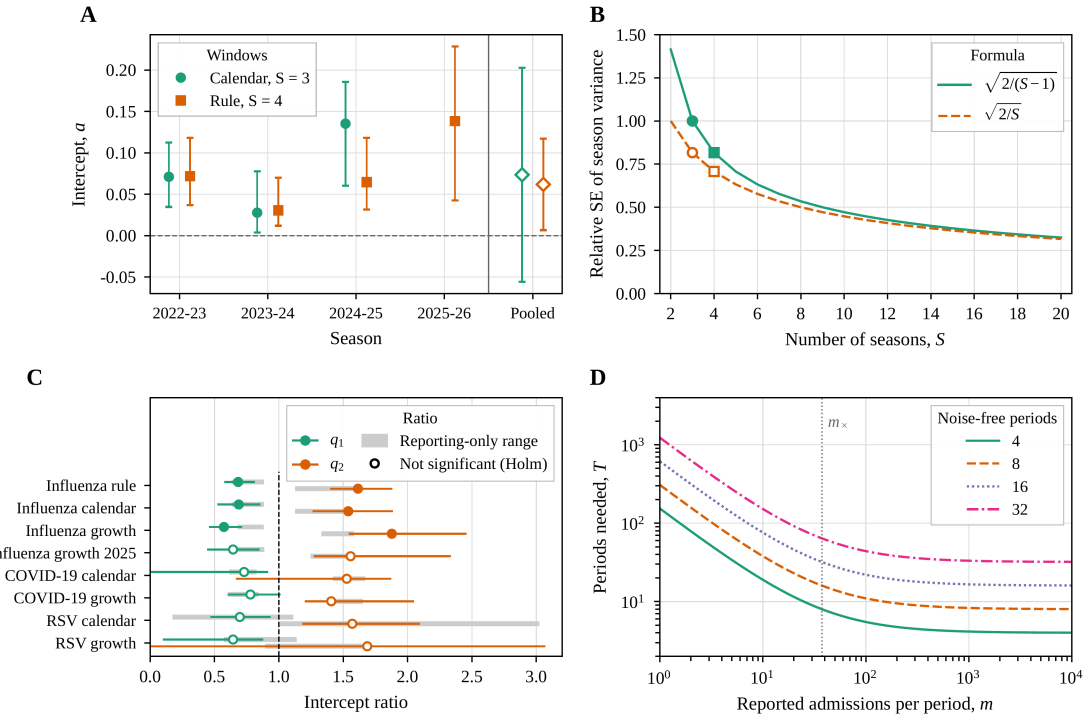


Fig. S6. Between-season uncertainty, intercept ratios across pathogens and number of periods. (A) Season-specific intercepts of the weighted growth ratio in calendar windows and in rule-defined early-growth windows, with two-way bootstrap 95% intervals; open diamonds show random-effects pooled estimates. Panel B gives the lower bound on the relative standard error of an estimated between-season variance against the number of seasons, compared with the values reported in the earlier analysis. Panel C gives intercept ratios for a single-week denominator and an equal three-week denominator relative to the original weights; grey bars give the ranges expected from reporting noise and epidemic growth without environmental variation, and filled markers denote one-sided tests that remain significant after Holm correction. Panel D gives the number of periods needed to match the precision of a given number of noise-free periods, by reported admissions per period; the panel is illustrative and is superseded by the empirical nomogram in Fig. S5A.

D. Supplementary tables

Tables S1–S7 report the estimates that underlie Figs. 3A, 4A and S1–S6. Table S8 gives the implied environmental variance under each setting of the inversion (Figs. 4B and S2), Table S9 lists the specifications used across the rounds of this analysis, and Table S11 gives the thinning and within-state results of Fig. 3B and C. Table S4 compares the voluntary reporting period with the same calendar months of the mandatory years. All values were generated from the stored outputs of the analysis code.

Table S1. Intercept a, slope b and balance ratio b/a of the relative loss by specification, with the two-way bootstrap 95% confidence interval of each estimate.

Specification	N	Weeks	$\bar{x}$	a [95% CI]	b [95% CI]	b/a [95% CI]	SD ratio
R_A, rule windows ($S = 4$)	1748	42	1.206	0.0150 [0.0100, 0.0208]	0.560 [0.367, 0.769]	37.4 [19.1, 71.0]	1.95
G, rule windows ($S = 4$)	1756	42	1.724	0.0840 [0.0498, 0.1248]	1.971 [0.959, 3.098]	23.5 [8.4, 59.7]	2.29
G, calendar windows ($S = 3$)	1267	29	1.671	0.0991 [0.0569, 0.1382]	2.243 [1.060, 3.559]	22.6 [9.0, 52.2]	4.68
R_A, calendar windows ($S = 3$)	1263	29	1.190	0.0202 [0.0097, 0.0306]	0.516 [0.279, 0.839]	25.6 [10.2, 79.6]	3.46

Iterated FGLS fit of $L = ((x - \bar{x})/\bar{x})^2$ on $a + b/D$; 95% percentile intervals from the two-way (state × week) pigeonhole bootstrap, $B = 2000$, $\bar{x}$ re-estimated in every replicate. SD ratio: bootstrap SD of a under two-way versus state-only resampling. 51 jurisdictions throughout.

Table S2. Decomposition of the intercept and comparison of the affine model with a power law, with the 95% confidence interval of each intercept.

Specification	a total	a after week effects	a after week and state-season effects	Week share	ICC	Power exponent	p	CV folds (state; season)
R_A, rule windows ($S = 4$)	0.0150 [0.0100, 0.0208]	0.0088 [0.0064, 0.0124]	0.0083 [0.0064, 0.0141]	0.41	0.25	0.38	0.11	26/51; 2/4
G, rule windows ($S = 4$)	0.0840 [0.0508, 0.1226]	0.0415 [0.0306, 0.0586]	0.0394 [0.0293, 0.0614]	0.51	0.31	0.35	0.44	19/51; 0/4
G, calendar windows ($S = 3$)	0.0991 [0.0583, 0.1399]	0.0473 [0.0337, 0.0676]	0.0450 [0.0335, 0.0932]	0.52	0.29	0.27	0.42	28/51; 3/3

Intervals are from a two-way bootstrap with $B = 1000$; the draws are separate from those of Table S1, which explains small differences. The third column removes national week effects from the statistic before the fit, and the fourth also removes state-season effects. The week share equals one minus the ratio of the third to the second column. The ICC is the share of the variance of the statistic that lies between weeks. The power exponent is c in $L = \alpha D^{-c}$, and p is from a parametric-bootstrap test of the affine model against this power law with 200 null draws. CV folds counts the leave-one-state-out and leave-one-season-out folds in which the affine model predicts better. A power analysis found no fraction of the estimated a (≤ 1) detectable with 80% power.

Table S3. Equivalence tests for the contribution of facility-reporting completeness to the intercept, with the paired bootstrap 90% confidence interval of each change.

Specification	$10^3 \Delta a$ [90% CI]	$10^3 \delta_{0.10}$	Equiv.	$10^3 \delta_{0.05}$	Equiv.	Seasons equiv.
R_A, rule windows ($S = 4$)	0.29 [−0.22, 0.85]	2.55	yes	0.64	no	4/4
G, rule windows ($S = 4$)	0.37 [−2.83, 2.81]	10.13	yes	2.53	no	3/4
G, calendar windows ($S = 3$)	−0.38 [−3.55, 2.24]	10.14	yes	2.54	no	2/3

Δa is the intercept of the statistic minus that of the statistic recomputed from counts divided by the share of hospitals reporting, with a paired two-way bootstrap ($B = 2000$). Margins $\delta_{\tau_0} = \kappa^2 \tau_0^2 (1 + v_R/\bar{R}^2)$ were set from a reporting-volatility threshold and the sample mean $\bar{R}$ of each statistic ($\kappa = 1$ for G, $v_R = 0.04$). Equivalence is declared when the 90% interval lies inside $(-\delta, \delta)$ (TOST, $\alpha = 0.05$). The column of equivalent seasons counts seasons equivalent at $\tau_0 = 0.10$.

Table S4. Calendar-matched comparison of the voluntary-reporting period (May–October).

Year	Reporting	N	SD of G	IQR of $\log G$
2023	mandatory	773	0.62 [0.40, 1.00]	0.48 [0.40, 0.61]
2024	voluntary	420	1.82 [0.42, 4.10]	0.76 [0.48, 1.04]
2025	mandatory	841	0.39 [0.33, 0.44]	0.48 [0.37, 0.58]

State-weeks with $D_t \geq 5$ whose week ends between 1 May and 31 October; 95% two-way bootstrap intervals, $B = 1000$. The FGLS intercept is not stable in the voluntary period (negative point estimate), so robust spread measures are reported. Laboratory proxy (Delphi FluView clinical laboratory data, 1224 state-weeks): share of the R_A intercept shared with laboratory growth 0.23 [0.043, 0.42]; correlation 0.31 [0.19, 0.41].

Table S5. Season-specific intercepts, each with its bootstrap 95% confidence interval, and between-season heterogeneity, reported separately for $S = 4$ and $S = 3$.

Season	R_A, rule windows ($S = 4$)	G, rule windows ($S = 4$)	G, calendar windows ($S = 3$)
2022-23	0.0137 [0.0066, 0.0213]	0.0718 [0.0368, 0.1181]	0.0711 [0.0347, 0.1124]
2023-24	0.0062 [0.0025, 0.0144]	0.0305 [0.0119, 0.0700]	0.0277 [0.0037, 0.0776]
2024-25	0.0127 [0.0057, 0.0241]	0.0647 [0.0315, 0.1181]	0.1352 [0.0602, 0.1857]
2025-26	0.0225 [0.0087, 0.0352]	0.1383 [0.0426, 0.2285]	–
a_{RE} [95% CI]	0.0125 [0.0022, 0.0227]	0.0619 [0.0066, 0.1171]	0.0735 [−0.0558, 0.2028]
Q (p)	5.71 (0.126)	6.58 (0.087)	8.92 (0.012)
$10^3\tau^2$ [Q-profile]	0.019 [0.000, 0.599]	0.620 [0.000, 26.756]	2.077 [0.127, 114.859]
I^2	0.47	0.54	0.78

Season intercepts use the season's own mean; 95% two-way bootstrap intervals, $B = 1000$. DerSimonian–Laird random effects with a t_{S-1} interval; Q-profile 95% interval for τ^2 (Viechtbauer, 2007). The relative standard error of a between-season variance estimated from S seasons is at least $\sqrt{2/(S-1)}$: 1.00 for $S = 3$ and 0.82 for $S = 4$.

Table S6. Cross-pathogen intercept ratios against the pure-reporting growth null, with Holm-adjusted p-values.

Pathogen, window	N	q_1 [95% CI]	q_2 [95% CI]	Null q_1	Null q_2	$q_1 < 1$	$q_2 > 1$	q_1 out	q_2 out	Excess > 0
Influenza, rule	1756	0.68 [0.58, 0.81]	1.61 [1.41, 1.87]	0.63–0.89	1.13–1.58	0.020	0.034	1.000	1.000	0.034
Influenza, calendar	1267	0.69 [0.53, 0.85]	1.54 [1.27, 1.88]	0.63–0.89	1.13–1.58	0.020	0.020	1.000	1.000	0.020
Influenza, growth	2707	0.57 [0.47, 0.71]	1.88 [1.55, 2.45]	0.68–0.88	1.33–1.58	0.034	0.020	1.000	1.000	0.020
Influenza, growth, 2025-26	663	0.64 [0.45, 0.84]	1.56 [1.28, 2.33]	0.66–0.89	1.25–1.57	0.116	0.075	1.000	1.000	0.046
COVID-19, calendar	1479	0.73 [n.d., 0.91]	1.53 [0.67, 1.87]	0.61–0.83	1.42–1.67	0.287	0.616	1.000	1.000	1.000
COVID-19, growth	2862	0.78 [0.61, 1.01]	1.41 [1.21, 2.04]	0.60–0.84	1.40–1.66	0.616	0.154	1.000	1.000	1.000
RSV, calendar	494	0.70 [0.48, 0.93]	1.57 [1.19, 2.09]	0.17–1.11	1.01–3.02	0.154	0.228	1.000	1.000	0.200
RSV, growth	1073	0.64 [0.11, 0.87]	1.69 [n.d., 3.06]	0.57–1.14	0.89–1.69	0.182	0.660	1.000	1.000	1.000

Here $q_1 = a_{\text{1 week}}/a_{\text{original}}$ and $q_2 = a_{\text{equal 3 weeks}}/a_{\text{original}}$, with intervals from a shared two-way bootstrap ($B = 2000$). Null ranges come from a growth-null simulation without transmission heterogeneity over three generation intervals, three admission delays and four reporting-noise levels; the rule-window row uses the influenza calendar range. The last five columns give one-sided bootstrap p-values, Holm-adjusted over all 40 tests, and the columns marked out test whether the ratio lies outside the null range. Excess denotes $a - v_R/\bar{G}^2$ with $v_R = 0.04$. A lower bound marked n.d. is not defined because some bootstrap replicates gave a negative intercept in the ratio, so the ratio has no positive lower limit; the upper limit is shown as estimated.

Table S7. Theory quantities recomputed from the closed forms and compared with an earlier version of this analysis.

Quantity	Earlier	Recomputed	Status
$m_\times$	53.1	53.08	reproduced
v_H ($\tau = 0.15$)	0.0307	0.0307	reproduced
$m_{\times,H}$ ($\tau = 0.15$)	30.4	30.36	reproduced
m_0^* ($c_0/c_1 = 100$)	72.9	72.86	reproduced
$T^\dagger$ ($\phi = 0.8$)	4.48	4.484	reproduced
$B^\dagger$	775	775.0	reproduced
$B^\dagger$ ($m_\times = 21.3$)	655	655.3	reproduced
$B^\dagger$ ($m_\times = 49.2$)	763	762.9	reproduced
Reporting lever term ($\times 10^{-3}$)	7.85	7.88	corrected
Lever total ($\times 10^{-3}$)	22.5	22.57	corrected
Fisher constant	0.505	0.505	reproduced
Relative SE, $S = 3$	0.82	1.000	corrected
Relative SE, $S = 4$	0.71	0.816	corrected

Baseline $\bar{R} = 1.15$, $v_R = 0.04$, $k = 0.35$, $\rho = 0.25$. The earlier relative standard errors for $S = 3, 4$ equal $\sqrt{2/S}$; the stated bound $\sqrt{2/(S-1)}$ gives the recomputed values. The lever terms differ by rounding.

Table S8. Implied environmental variance v_R^{impl} from the two-step inversion under each simulation setting, with the 95% confidence interval of each estimate.

Setting	v_R^{impl} from R_A	v_R^{impl} from G	κ_{R_A}	κ_G	Ratio R_A/G	Agree
Data-matched design (M0)	0.040 [0.026, 0.056]	0.043 [0.025, 0.064]	0.516	2.729	0.93 [0.78, 1.18]	yes
No national shock	0.039 [0.026, 0.055]	0.042 [0.024, 0.062]	0.523	2.775	0.94 [0.79, 1.18]	yes
National share 0.65	0.040 [0.026, 0.056]	0.043 [0.025, 0.065]	0.510	2.670	0.92 [0.77, 1.16]	yes
AR(1), $\phi = 0.5$	0.027 [0.017, 0.040]	0.025 [0.013, 0.039]	0.780	4.699	1.08 [0.88, 1.43]	yes
True GI mean 2.2 d	0.013 [0.008, 0.019]	0.014 [0.008, 0.021]	1.287	6.991	0.94 [0.77, 1.20]	yes
True GI mean 4.3 d	0.058 [0.038, 0.081]	0.063 [0.037, 0.095]	0.380	1.977	0.92 [0.77, 1.16]	yes
Delay mean 3.5 d	0.038 [0.025, 0.053]	0.041 [0.024, 0.061]	0.535	2.790	0.93 [0.77, 1.16]	yes
Delay mean 8 d	0.046 [0.030, 0.064]	0.047 [0.027, 0.069]	0.454	2.483	0.98 [0.82, 1.22]	yes
$k = 0.1$	0.040 [0.026, 0.055]	0.042 [0.025, 0.062]	0.517	2.740	0.94 [0.79, 1.18]	yes
Synchronous count shock	0.019 [0.006, 0.035]	0.029 [0.012, 0.048]	1.041	4.553	0.67 [0.39, 0.90]	no
Generic design (bridge)	0.043 [0.029, 0.059]	0.044 [0.027, 0.063]	0.507	2.734	0.97 [0.83, 1.18]	–

Each setting changes one element of the data-matched design (Section 3.10); the analyst's generation interval stays at a mean of 3.6 days throughout. Intervals combine 2000 joint two-way bootstrap draws of the two empirical intercepts with draws of the calibration line from its Monte Carlo sampling distribution. $\kappa = \mathbb{E}(\hat{a})/(v_R/\bar{R}^2)$ at the grid value of v_R nearest to the implied variance. Agree marks settings in which the 95% interval of the difference between the two implied variances contains zero; the generic design of the first simulation is shown as a bridge and is not counted among the ten settings.

Table S9. Specifications used across the rounds of this analysis, with the 95% confidence interval of each estimate on the data version of 2026-10-20.

Round and specification	Role	a [95% CI]	b/a [95% CI]
Round 1: G, three calendar windows ($S = 3$)	Primary in round 1; secondary from round 2	0.0991 [0.0569, 0.1382]	22.6 [9.0, 52.2]
Round 2, rule version 1: windows ending at the season maximum	Abandoned (four weeks in 2024–25); no estimates retained	–	–
Round 2, rule version 2 (2026-10-20): G, rule windows ($S = 4$)	Secondary	0.0840 [0.0498, 0.1248]	23.5 [8.4, 59.7]
Round 2, rule version 2: R_A, rule windows ($S = 4$)	Primary from round 2	0.0150 [0.0100, 0.0208]	37.4 [19.1, 71.0]
Round 2: R_A, calendar windows ($S = 3$)	Sensitivity	0.0202 [0.0097, 0.0306]	25.6 [10.2, 79.6]
Round 3: two-step inversion	Calibration of the round-2 fits; no refit	–	–

All estimates were refitted on the same data version with two-way bootstrap intervals ($B = 2000$). The renewal-type ratio became the primary statistic in round 2 because its relationship to R_t is explicit through the generation interval; the weighted growth ratio is retained for continuity. The point estimate of each specification lies inside the interval of the other specification of the same statistic.

Table S10. Lag-1 autocorrelation, and implied environmental variance and b/a in the added inversion settings.

A. Lag-1 autocorrelation

Setting	r_1 of R_A	Data pct.	r_1 of G	Data pct.
Data (rule windows)	−0.095 [−0.168, −0.019]	–	0.123 [0.039, 0.201]	–
Data-matched design (M0)	−0.074 [−0.151, 0.014]	31.0	0.114 [0.025, 0.220]	62.0
AR(1), $\phi = 0.25$	0.007 [−0.098, 0.118]	3.0	0.193 [0.097, 0.304]	10.0
AR(1), $\phi = 0.5$	0.072 [−0.039, 0.229]	0.0	0.258 [0.146, 0.390]	0.0
True GI mean 2.2 d	0.099 [−0.002, 0.223]	0.0	0.251 [0.156, 0.363]	0.0
True GI mean 2.9 d	0.006 [−0.083, 0.120]	2.5	0.183 [0.091, 0.296]	9.0
True GI mean 4.3 d	−0.134 [−0.206, −0.030]	80.0	0.051 [−0.031, 0.161]	92.0
Admissions VMR 1.5	−0.170 [−0.233, −0.096]	97.5	−0.014 [−0.092, 0.076]	100.0
Admissions VMR 2.0	−0.183 [−0.247, −0.104]	99.0	−0.041 [−0.104, 0.035]	100.0
Synchronous count shock	−0.146 [−0.225, −0.076]	86.0	0.041 [−0.033, 0.116]	99.0

B. Added settings (b/a pooled at $v_R = 0.04$)

Setting	v_R^{impl} from R_A	v_R^{impl} from G	b/a of R_A	b/a of G
Data-matched design (M0)	0.040 [0.026, 0.056]	0.043 [0.025, 0.064]	28.0	16.2
True GI mean 2.9 d	0.026 [0.017, 0.036]	0.027 [0.016, 0.040]	18.2	10.4
Admissions VMR 1.5	0.039 [0.026, 0.055]	0.044 [0.026, 0.066]	68.0	40.2
Admissions VMR 2.0	0.039 [0.026, 0.055]	0.045 [0.027, 0.066]	81.0	48.4
AR(1), $\phi = 0.25$	0.033 [0.021, 0.047]	0.033 [0.019, 0.049]	23.4	12.6

r_1 is the pooled lag-1 autocorrelation of the relative deviation $x/\bar{x}_{\text{block}} - 1$ within state-season blocks. The interval for the data comes from a state-cluster bootstrap with 1000 draws; simulated rows give the mean and the 2.5% to 97.5% range of r_1 over panels at $v_R = 0.04$ (200 panels per value of v_R for the added settings, 100 for the others), and Data pct. is the percentile of the empirical value in that distribution. The analyst's generation interval stays at a mean of 3.6 days; 2.9 and 4.3 days are the ends of the 95% interval for the serial interval of influenza in Cowling et al. (2009). Overdispersed admissions use a gamma–Poisson mixture on weekly counts.

Table S11. Binomial thinning of the renewal-type ratio sample and fits with state intercepts. Panel A gives, for each retention probability q, the slope on the original scale C_t, the implied thinning slope $c_q = [b(q) - b(1)]/(1/q - 1)$, the change in intercept $a(q) - a(1)$ and c_q in the subset with $C_t \geq 50$ ($n = 828$). Panel B gives the slope on $1/C_t$ and the ratio of slope to intercept under each specification. Values are point estimates with 95% two-way bootstrap intervals in brackets. The intervals in Panel B come from a single set of 1000 draws shared by all specifications, so the pooled interval differs slightly from that in Table S1.

A. Binomial thinning ($n = 1748$)

q	$b(q)$	c_q	Δa ($\times 10^{-3}$)	c_q, $C_t \geq 50$
0.8	0.678 [0.447, 0.918]	0.474 [0.038, 0.956]	−0.34 [−1.84, 1.17]	0.472 [−0.509, 1.416]
0.6	0.881 [0.609, 1.197]	0.481 [0.189, 0.883]	−0.89 [−4.13, 1.46]	0.441 [−0.123, 1.071]
0.5	1.053 [0.728, 1.397]	0.493 [0.230, 0.818]	−1.30 [−4.78, 1.84]	0.458 [−0.032, 0.988]
0.4	1.316 [0.917, 1.788]	0.504 [0.264, 0.785]	−1.90 [−6.26, 1.87]	0.451 [0.074, 0.946]
0.3	1.731 [1.256, 2.371]	0.502 [0.319, 0.765]	−2.40 [−7.48, 1.54]	0.458 [0.086, 0.827]
0.2	2.457 [1.736, 3.181]	0.474 [0.307, 0.652]	−2.50 [−8.50, 3.39]	0.457 [0.195, 0.773]
0.1	4.010 [2.913, 5.346]	0.383 [0.265, 0.532]	0.44 [−7.01, 11.93]	0.483 [0.241, 0.857]
Pooled slope c	–	0.407 [0.311, 0.531]	–	0.477 [0.286, 0.769]
Poisson coefficient b_P	–	0.446	–	0.420

B. Within-state fixed effects ($n = 1748$, 51 states)

Specification	Slope	Ratio
Pooled cross-section	0.560 [0.368, 0.771]	37.4 [18.6, 73.9]
State intercepts, mean a_s	0.379 [0.104, 0.669]	17.3 [3.3, 51.4]
State intercepts, median a_s	–	21.7 [3.7, 60.6]
State intercepts, row-weighted a_s	–	18.4 [3.5, 49.9]
State intercepts, case-weighted a_s	–	24.6 [4.9, 58.4]
State intercepts, scale D_t	0.369 [0.111, 0.643]	17.7 [3.2, 51.4]
Week intercepts only	0.508 [0.232, 0.835]	–
State and week intercepts	0.78 [−1.20, 2.77]	–
Paired ratio, cross-section to mean a_s	–	2.16 [0.93, 7.75]

References

Marie Auger-Méthé, Chris Field, Christoffer M. Albertsen, Andrew E. Derocher, Mark A. Lewis, Ian D. Jonsen, and Joanna Mills Flemming. State-space models' dirty little secrets: even simple linear Gaussian models can have estimation problems. *Scientific Reports*, 6:26677, 2016.

Amin Azmon, Christel Faes, and Niel Hens. On the estimation of the reproduction number based on misreported epidemic data. *Statistics in Medicine*, 33(7):1176–1192, 2014.

Yoav Benjamini and Yosef Hochberg. Controlling the false discovery rate: A practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society: Series B (Methodological)*, 57(1):289–300, 1995.

Samir Bhatt, Neil Ferguson, Seth Flaxman, Axel Gandy, Swapnil Mishra, and James A. Scott. Semi-mechanistic Bayesian modelling of COVID-19 with renewal processes. *Journal of the Royal Statistical Society Series A: Statistics in Society*, 186(4):601–615, 2023.

Carles Bretó, Daihai He, Edward L. Ionides, and Aaron A. King. Time series analysis via mechanistic models. *The Annals of Applied Statistics*, 3(1):319–348, 2009.

A. Colin Cameron, Jonah B. Gelbach, and Douglas L. Miller. Bootstrap-based improvements for inference with clustered errors. *The Review of Economics and Statistics*, 90(3):414–427, 2008.

A. Colin Cameron, Jonah B. Gelbach, and Douglas L. Miller. Robust inference with multiway clustering. *Journal of Business & Economic Statistics*, 29(2):238–249, 2011.

Simon Cauchemez, Alain-Jacques Valleron, Pierre-Yves Boëlle, Antoine Flahault, and Neil M. Ferguson. Estimating the impact of school closure on influenza transmission from Sentinel data. *Nature*, 452(7188):750–754, 2008.

Centers for Disease Control and Prevention. FluSight: Flu forecasts, May 3, 2022, 2022.

Centers for Disease Control and Prevention. Weekly hospital respiratory data (HRD) metrics by jurisdiction, national healthcare safety network (NHSN), 2026.

Centers for Medicare and Medicaid Services. QSO-25-05: Hospital respiratory data reporting requirements, 2024.

Louis Yat Hin Chan, Sinead E. Morris, Melissa S. Stockwell, Natalie M. Bowman, Edwin Asturias, Suchitra Rao, Karen Lutrick, Katherine D. Ellingson, Huong Q. Nguyen, Yvonne Maldonado, Son H. McLaren, Ellen Sano, Jessica E. Biddle, Sarah E. Smith-Jeffcoat, Matthew Biggerstaff, Melissa A. Rolfes, H. Keipp Talbot, Carlos G. Grijalva, Rebecca K. Borchering, and Alexandra M. Mellis. Estimating the generation time for influenza transmission using household data in the United States. *Epidemics*, 50:100815, 2025.

William G. Cochran. *Sampling Techniques*. John Wiley & Sons, New York, 3rd edition, 1977.

Anne Cori, Neil M. Ferguson, Christophe Fraser, and Simon Cauchemez. A new framework and software to estimate time-varying reproduction numbers during epidemics. *American Journal of Epidemiology*, 178(9):1505–1512, 2013.

Benjamin J. Cowling, Vicky J. Fang, Steven Riley, J. S. Malik Peiris, and Gabriel M. Leung. Estimation of the serial interval of influenza. *Epidemiology*, 20(3):344–347, 2009.

Benjamin D. Dalziel, Stephen Kissler, Julia R. Gog, Cécile Viboud, Ottar N. Bjørnstad, C. Jessica E. Metcalf, and Bryan T. Grenfell. Urbanization and humidity shape the intensity of influenza epidemics in U.S. cities. *Science*, 362(6410):75–79, 2018.

Rebecca DerSimonian and Nan Laird. Meta-analysis in clinical trials. *Controlled Clinical Trials*, 7(3):177–188, 1986.

J. Dureau, K. Kalogeropoulos, and M. Baguelin. Capturing the time-varying drivers of an epidemic using stochastic dynamical systems. *Biostatistics*, 14(3):541–555, 2013.

Oliver Eales and Steven Riley. Differences between the true reproduction number and the apparent reproduction number of an epidemic time series. *Epidemics*, 46:100742, 2024.

Akira Endo, Centre for the Mathematical Modelling of Infectious Diseases COVID-19 Working Group, Sam Abbott, Adam J. Kucharski, and Sebastian Funk. Estimating the overdispersion in COVID-19 transmission using outbreak sizes outside China. *Wellcome Open Research*, 5:67, 2020.

Steinar Engen, Øyvind Bakke, and Aminul Islam. Demographic and environmental stochasticity—concepts and definitions. *Biometrics*, 54(3):840–846, 1998.

Seth Flaxman, Swapnil Mishra, Axel Gandy, H. Juliette T. Unwin, Thomas A. Mellan, Helen Coupland, Charles Whittaker, Harrison Zhu, Tresnia Berah, Jeffrey W. Eaton, Mélodie Monod, et al. Estimating the effects of non-pharmaceutical interventions on COVID-19 in Europe. *Nature*, 584(7820):257–261, 2020.

Christophe Fraser. Estimating individual and household reproduction numbers in an emerging epidemic. *PLoS ONE*, 2 (8):e758, 2007.

Aaron M. Frutos, Shunte Moon, Alison M. Binder, Andrea J. Cool, Kingsley Iyawe, Trevor Thompson, Margaret A. Dudeck, Catherine H. Bozio, Mark W. Tenforde, Matthew Biggerstaff, and Kelsey M. Sumner. Evaluation of a novel data source for national influenza surveillance: Influenza hospitalization data in the National Healthcare Safety Network, United States, September 2021–April 2024. *American Journal of Public Health*, 116(2):251–255, 2026.

Katelyn M. Gostic, Lauren McGough, Edward B. Baskerville, Sam Abbott, Keya Joshi, Christine Tedijanto, Rebecca Kahn, Rene Niehus, James A. Hay, Pablo M. De Salazar, Joel Hellewell, Sophie Meakin, James D. Munday, Nikos I. Bosse, Katharine Sherratt, Robin N. Thompson, Laura F. White, Jana S. Huisman, Jérémie Scire, Sebastian Bonhoeffer, Tanja Stadler, Jacco Wallinga, Sebastian Funk, Marc Lipsitch, and Sarah Cobey. Practical considerations for measuring the effective reproductive number, R_t. *PLOS Computational Biology*, 16(12):e1008409, 2020.

James D. Hamilton. *Time Series Analysis*. Princeton University Press, Princeton, NJ, 1994.

R. J. Hayes and S. Bennett. Simple sample size calculation for cluster-randomized trials. *International Journal of Epidemiology*, 28(2):319–326, 1999.

Daihai He, Edward L. Ionides, and Aaron A. King. Plug-and-play inference for disease dynamics: measles in large and small populations as a case study. *Journal of the Royal Society Interface*, 7(43):271–283, 2010.

Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979.

Jana S. Huisman, Jérémie Scire, Daniel C. Angst, Jinzhou Li, Richard A. Neher, Marloes H. Maathuis, Sebastian Bonhoeffer, and Tanja Stadler. Estimation and worldwide monitoring of the effective reproductive number of SARS-CoV-2. *eLife*, 11:e71345, 2022.

Samuel Karlin and Herman Rubin. The theory of decision procedures for distributions with monotone likelihood ratio. *The Annals of Mathematical Statistics*, 27(2):272–299, 1956.

Thomas M. Kincaid, David P. Larsen, and N. Scott Urquhart. The structure of variation and its influence on the estimation of status: Indicators of condition of lakes in the Northeast, U.S.A. *Environmental Monitoring and Assessment*, 98(1–3): 1–21, 2004.

Aaron A. King, Matthieu Domenech de Cellès, Felicia M. G. Magpantay, and Pejman Rohani. Avoidable errors in the modelling of outbreaks of emerging pathogens, with special reference to Ebola. *Proceedings of the Royal Society B: Biological Sciences*, 282(1806):20150347, 2015.

Jonas Knape. Estimability of density dependence in models of time series data. *Ecology*, 89(11):2994–3000, 2008.

Adam J. Kucharski, Timothy W. Russell, Charlie Diamond, Yang Liu, John Edmunds, Sebastian Funk, Rosalind M. Eggo, et al. Early dynamics of transmission and control of COVID-19: a mathematical modelling study. *The Lancet Infectious Diseases*, 20(5):553–558, 2020.

Russell Lande. Risks of population extinction from demographic and environmental stochasticity and random catastrophes. *The American Naturalist*, 142(6):911–927, 1993.

Russell Lande, Steinar Engen, and Bernt-Erik Sæther. *Stochastic Population Dynamics in Ecology and Conservation*. Oxford University Press, Oxford, 2003.

Erich L. Lehmann and George Casella. *Theory of Point Estimation*. Springer, New York, 2nd edition, 1998.

Erich L. Lehmann and Joseph P. Romano. *Testing Statistical Hypotheses*. Springer, New York, 3rd edition, 2005.

J. O. Lloyd-Smith, S. J. Schreiber, P. E. Kopp, and W. M. Getz. Superspreading and the effect of individual variation on disease emergence. *Nature*, 438(7066):355–359, 2005.

Sarabeth M. Mathis, Alexander E. Webber, Tomás M. León, Erin L. Murray, Monica Sun, Lauren A. White, Logan C. Brooks, Alden Green, Addison J. Hu, Roni Rosenfeld, et al. Evaluation of FluSight influenza forecasting in the 2021–22 and 2022–23 seasons with a new target laboratory-confirmed influenza hospitalizations. *Nature Communications*, 15:6289, 2024.

Peter McCullagh. Resampling and exchangeable arrays. *Bernoulli*, 6(2):285–301, 2000.

Rebecca K. Nash, Samir Bhatt, Anne Cori, and Pierre Nouvellet. Estimating the epidemic reproduction number from temporally aggregated incidence data: A statistical modelling approach and software tool. *PLOS Computational Biology*, 19(8):e1011439, 2023.

Pierre Nouvellet, Anne Cori, Tini Garske, Isobel M. Blake, Ilaria Dorigatti, Wes Hinsley, Thibaut Jombart, Harriet L. Mills, Gemma Nedjati-Gilani, Maria D. Van Kerkhove, Christophe Fraser, Christl A. Donnelly, Neil M. Ferguson, and Steven Riley. A simple approach to measure transmissibility and forecast incidence. *Epidemics*, 22:29–35, 2018.

I. Ogi-Gittins, W. S. Hart, J. Song, R. K. Nash, J. Polonsky, A. Cori, E. M. Hill, and R. N. Thompson. A simulation-based approach for estimating the time-dependent reproduction number from temporally aggregated disease incidence time series data. *Epidemics*, 47:100773, 2024.

Art B. Owen. The pigeonhole bootstrap. *The Annals of Applied Statistics*, 1(2):386–411, 2007.

Kris V. Parag. Improved estimation of time-varying reproduction numbers at low case incidence and between epidemic waves. *PLOS Computational Biology*, 17(9):e1009347, 2021.

Kris V. Parag and Christl A. Donnelly. Fundamental limits on inferring epidemic resurgence in real time using effective reproduction numbers. *PLOS Computational Biology*, 18(4):e1010004, 2022.

Kris V. Parag, Christl A. Donnelly, and Alexander E. Zarebski. Quantifying the information in noisy epidemic curves. *Nature Computational Science*, 2(9):584–594, 2022.

Kris V. Parag, Ben Lambert, Christl A. Donnelly, and Sandor Beregi. Asymmetric limits on timely interventions from noisy epidemic data. *Communications Physics*, 8:450, 2025.

Stephen W. Raudenbush. Statistical analysis and optimal design for cluster randomized trials. *Psychological Methods*, 2 (2):173–185, 1997.

Pejman Rohani, Matthew J. Keeling, and Bryan T. Grenfell. The interplay between determinism and stochasticity in childhood diseases. *The American Naturalist*, 159(5):469–481, 2002.

Donald J. Schuirmann. A comparison of the Two One-Sided Tests Procedure and the Power Approach for assessing the equivalence of average bioavailability. *Journal of Pharmacokinetics and Biopharmaceutics*, 15(6):657–680, 1987.

Shayle R. Searle, George Casella, and Charles E. McCulloch. *Variance Components*. John Wiley & Sons, New York, 1992.

Jeffrey Shaman and Melvin Kohn. Absolute humidity modulates influenza survival, transmission, and seasonality. *Proceedings of the National Academy of Sciences*, 106(9):3243–3248, 2009.

Jeffrey Shaman, Virginia E. Pitzer, Cécile Viboud, Bryan T. Grenfell, and Marc Lipsitch. Absolute humidity and the seasonal onset of influenza in the continental United States. *PLoS Biology*, 8(2):e1000316, 2010.

Nicholas Steyn, Kris V. Parag, Robin N. Thompson, and Christl A. Donnelly. A primer on inference and prediction with epidemic renewal models and sequential Monte Carlo. *Statistics in Medicine*, 44(18-19):e70204, 2025.

Oliver Stoner, Theo Economou, and Gabriela Drummond Marques da Silva. A hierarchical framework for correcting under-reporting in count data. *Journal of the American Statistical Association*, 114(528):1481–1492, 2019.

R. N. Thompson, J. E. Stockwin, R. D. van Gaalen, J. A. Polonsky, Z. N. Kamvar, P. A. Demarsh, E. Dahlqwist, S. Li, E. Miguel, T. Jombart, J. Lessler, S. Cauchemez, and A. Cori. Improved inference of time-varying reproduction numbers during infectious disease outbreaks. *Epidemics*, 29:100356, 2019.

Jerome I. Tokars, Sonja J. Olsen, and Carrie Reed. Seasonal incidence of symptomatic influenza in the United States. *Clinical Infectious Diseases*, 66(10):1511–1518, 2018.

N. Scott Urquhart and Thomas M. Kincaid. Designs for detecting trend from repeated surveys of ecological resources. *Journal of Agricultural, Biological, and Environmental Statistics*, 4(4):404–414, 1999.

Wolfgang Viechtbauer. Confidence intervals for the amount of heterogeneity in meta-analysis. *Statistics in Medicine*, 26 (1):37–52, 2007.

Margaretha Annelie Vink, Martinus Christoffel Jozef Bootsma, and Jacco Wallinga. Serial intervals of respiratory infectious diseases: A systematic review and analysis. *American Journal of Epidemiology*, 180(9):865–875, 2014.

J. Wallinga and M. Lipsitch. How generation intervals shape the relationship between growth rates and reproductive numbers. *Proceedings of the Royal Society B: Biological Sciences*, 274(1609):599–604, 2007.

Jacco Wallinga and Peter Teunis. Different epidemic curves for severe acute respiratory syndrome reveal similar impacts of control measures. *American Journal of Epidemiology*, 160(6):509–516, 2004.

Jeffrey M. Wooldridge. *Econometric Analysis of Cross Section and Panel Data*. MIT Press, Cambridge, MA, 2nd edition, 2010.